

Distributed Batch Analytics with MapReduce

Prof. Pamod Bhatotia

<https://dse.in.tum.de/bhatotia/>



How much data?



processes 20 PB a day (2008)
crawls 20B web pages a day (2012)



>10 PB data, 75B DB
calls per day (6/2012)

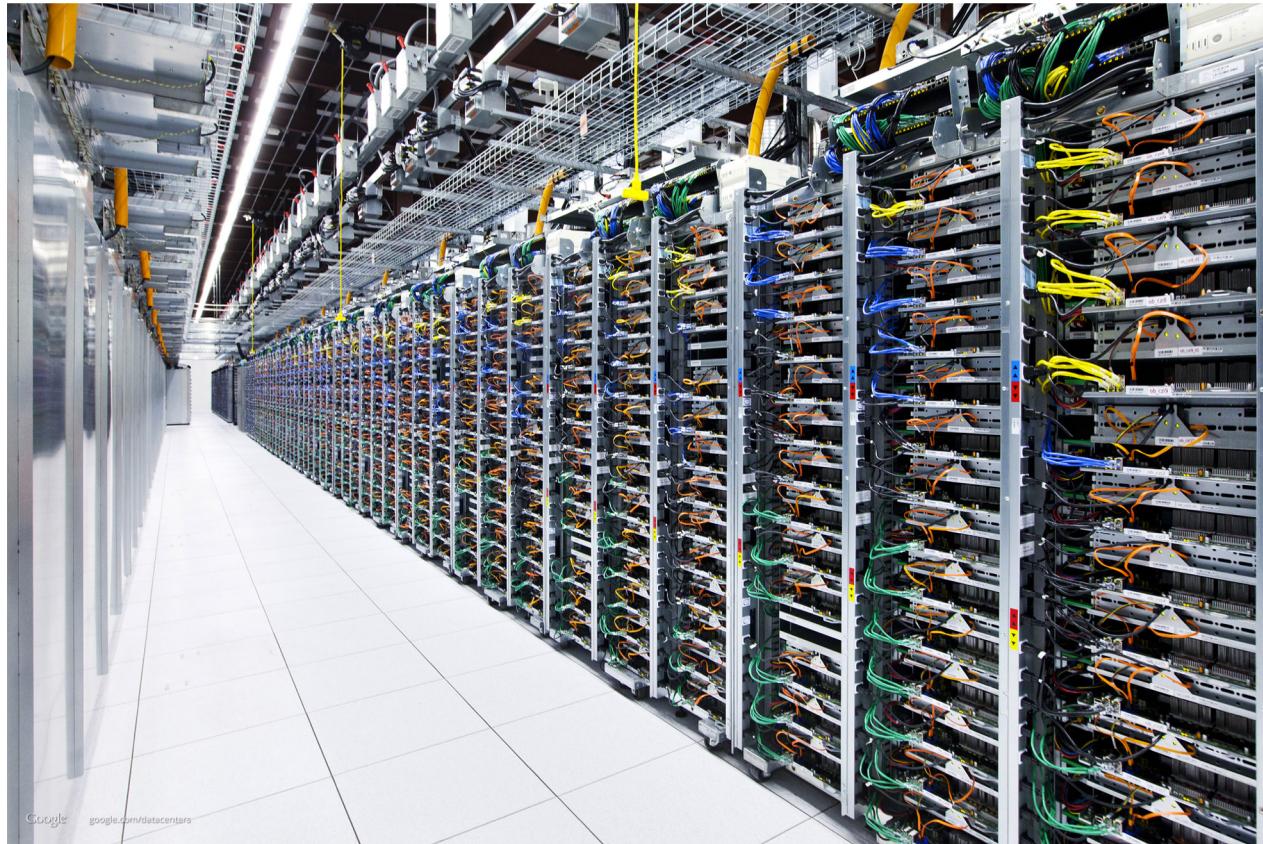
>100 PB of user data +
500 TB/day (8/2012)



S3: 449B objects, peak 290k
request/second (7/2011)
1T objects (6/2012)

Distributed Systems!

Data-center



Cluster of 100s of thousands of machines

In today's class

How to easily write parallel applications on distributed computing systems?

1. MapReduce
2. Pig
 - A high-level language built on top of MapReduce

Design challenges

- How to parallelize application logic?
- How to communicate?
- How to synchronize?
- How to perform load balancing?
- How to handle faults?
- How to schedule jobs?

For each and every application!

Design

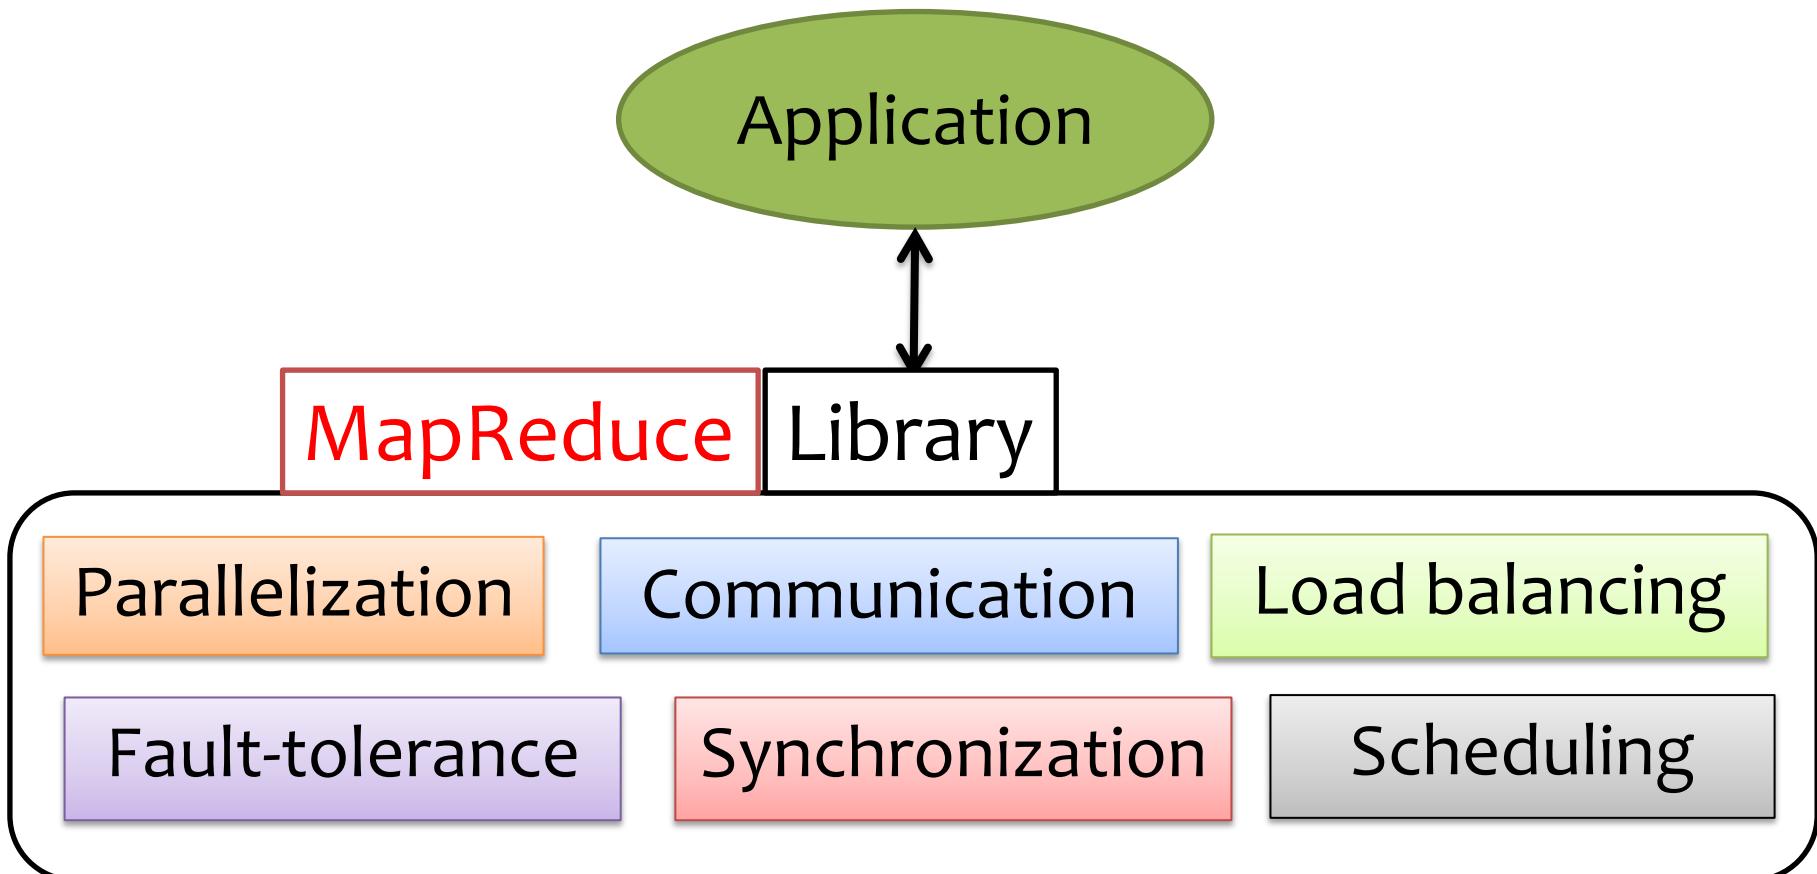
Implement

Optimize

Debug

Maintain

The power of abstraction



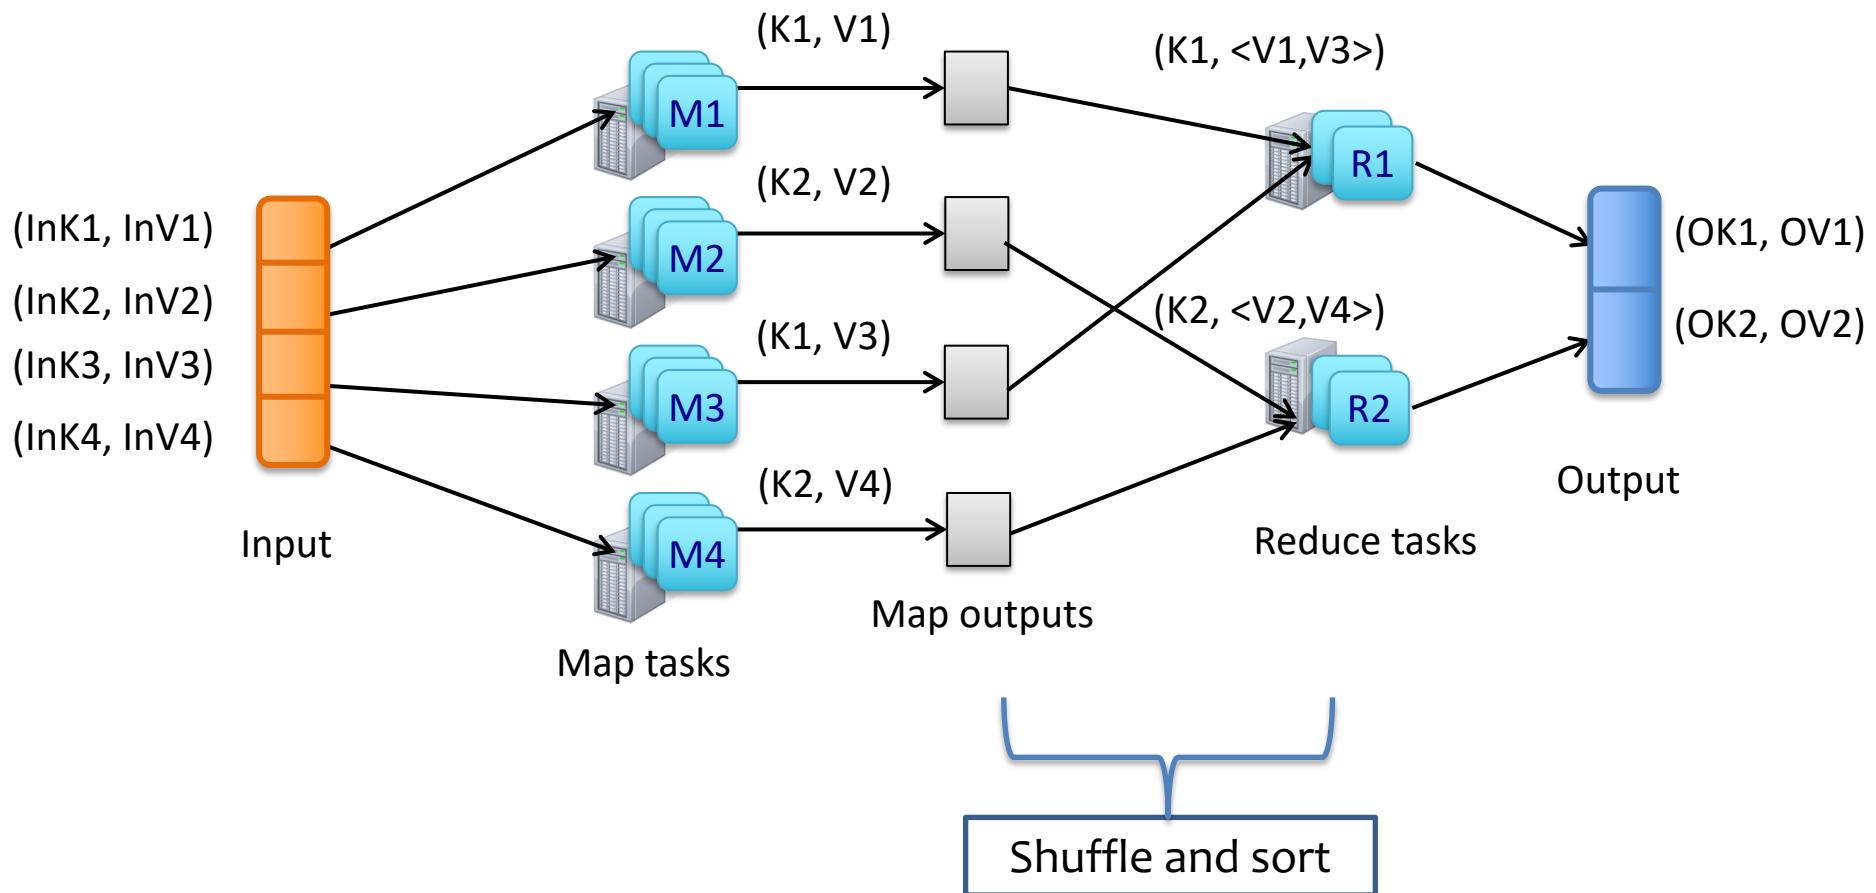
MapReduce

- Programming model
 - Programmer writes two methods: Map & Reduce
- Run-time library
 - Takes care of everything else!

MapReduce programming model

- Inspired from functional programming
 - Data-parallel application logic
- Programmer's interface:
 - $\text{Map}(\text{key}, \text{value}) \rightarrow (\text{key}, \text{value})$
 - $\text{Reduce}(\text{key}, \langle \text{value} \rangle) \rightarrow (\text{key}, \text{value})$

MapReduce run-time system



An example: word-count

- **Input:**
 - Given a corpus of documents, such as Wikipedia
- **Output:**
 - Count the frequency of each distinct word

MapReduce for word-count

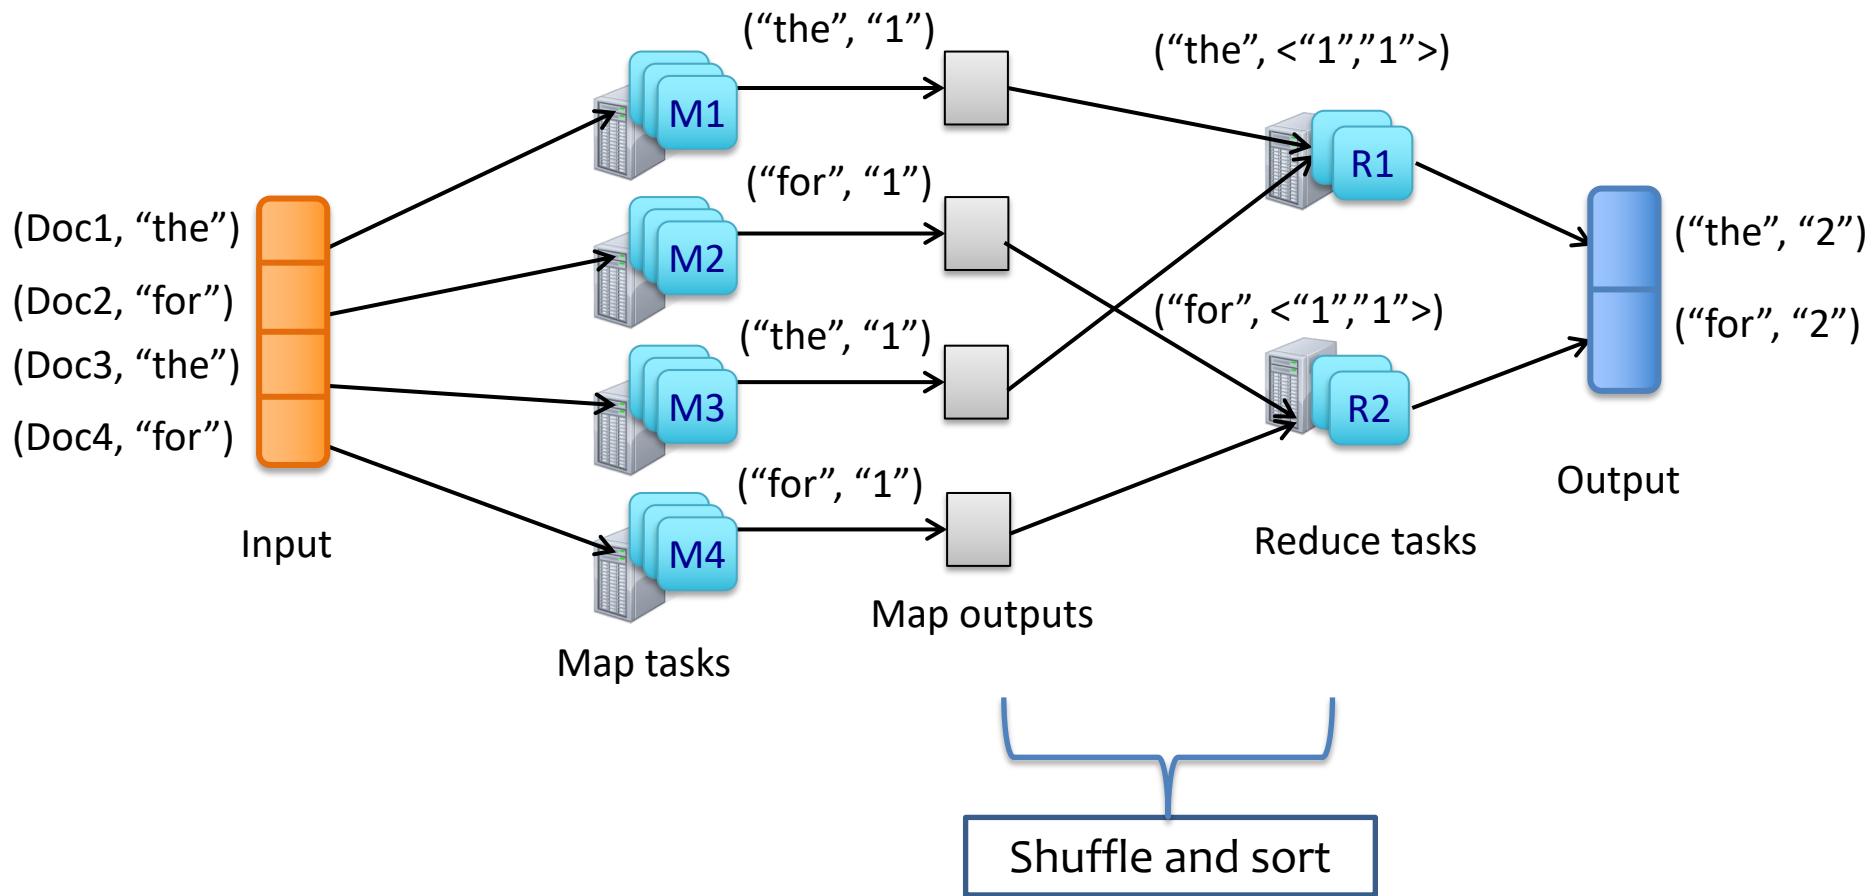
- **map(string key, string value)**

```
//key: document name  
//value: document contents  
for each word w in value  
    EmitIntermediate(w, "1");
```

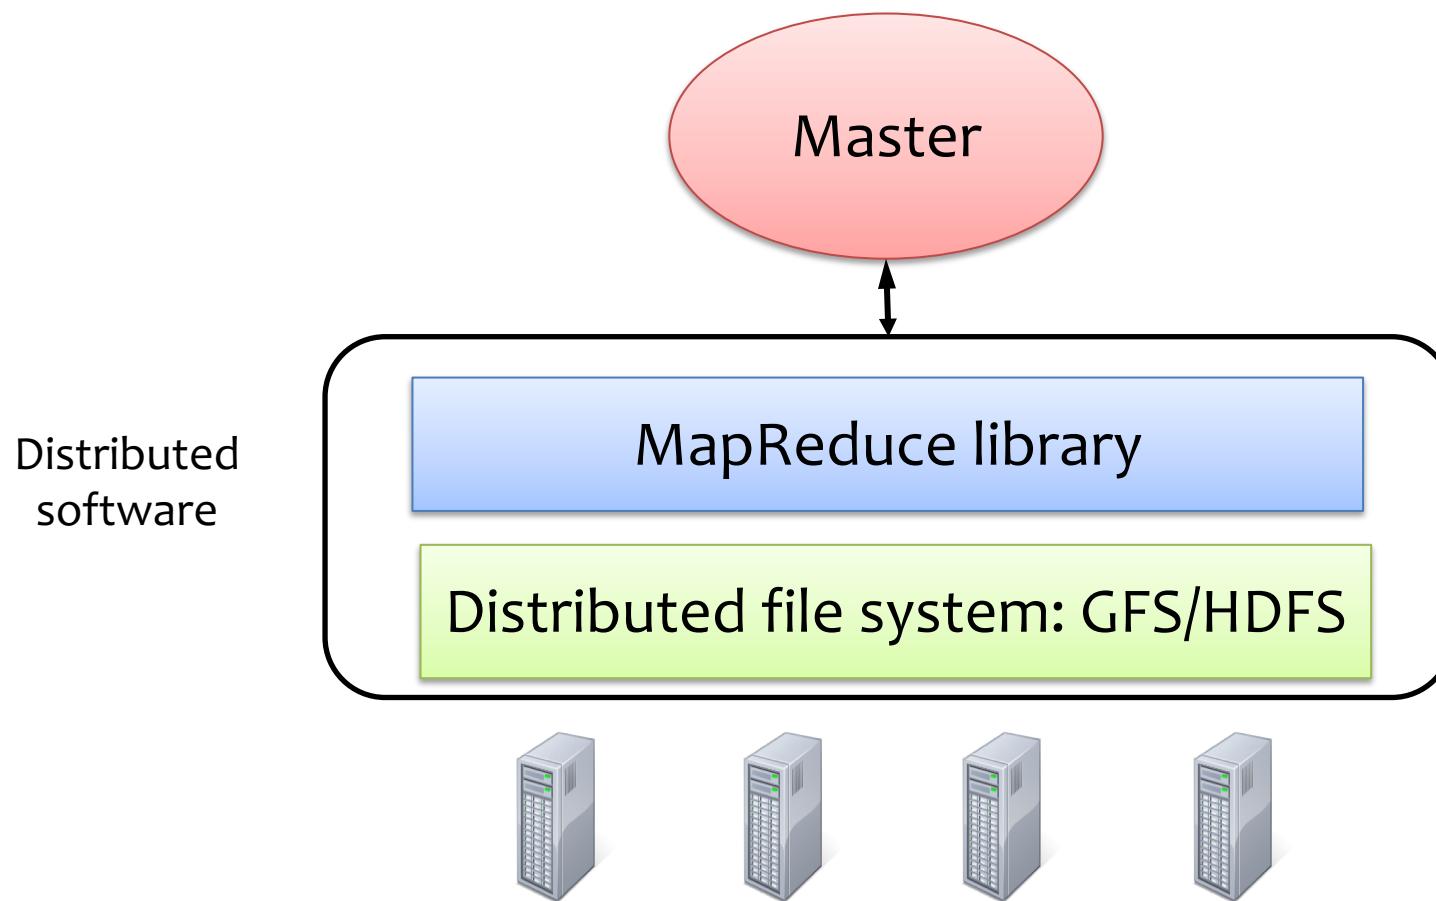
- **reduce(string key, iterator values)**

```
//key: word  
//values: list of counts  
int results = 0;  
for each v in values  
    result += ParseInt(v);  
Emit(key, AsString(result));
```

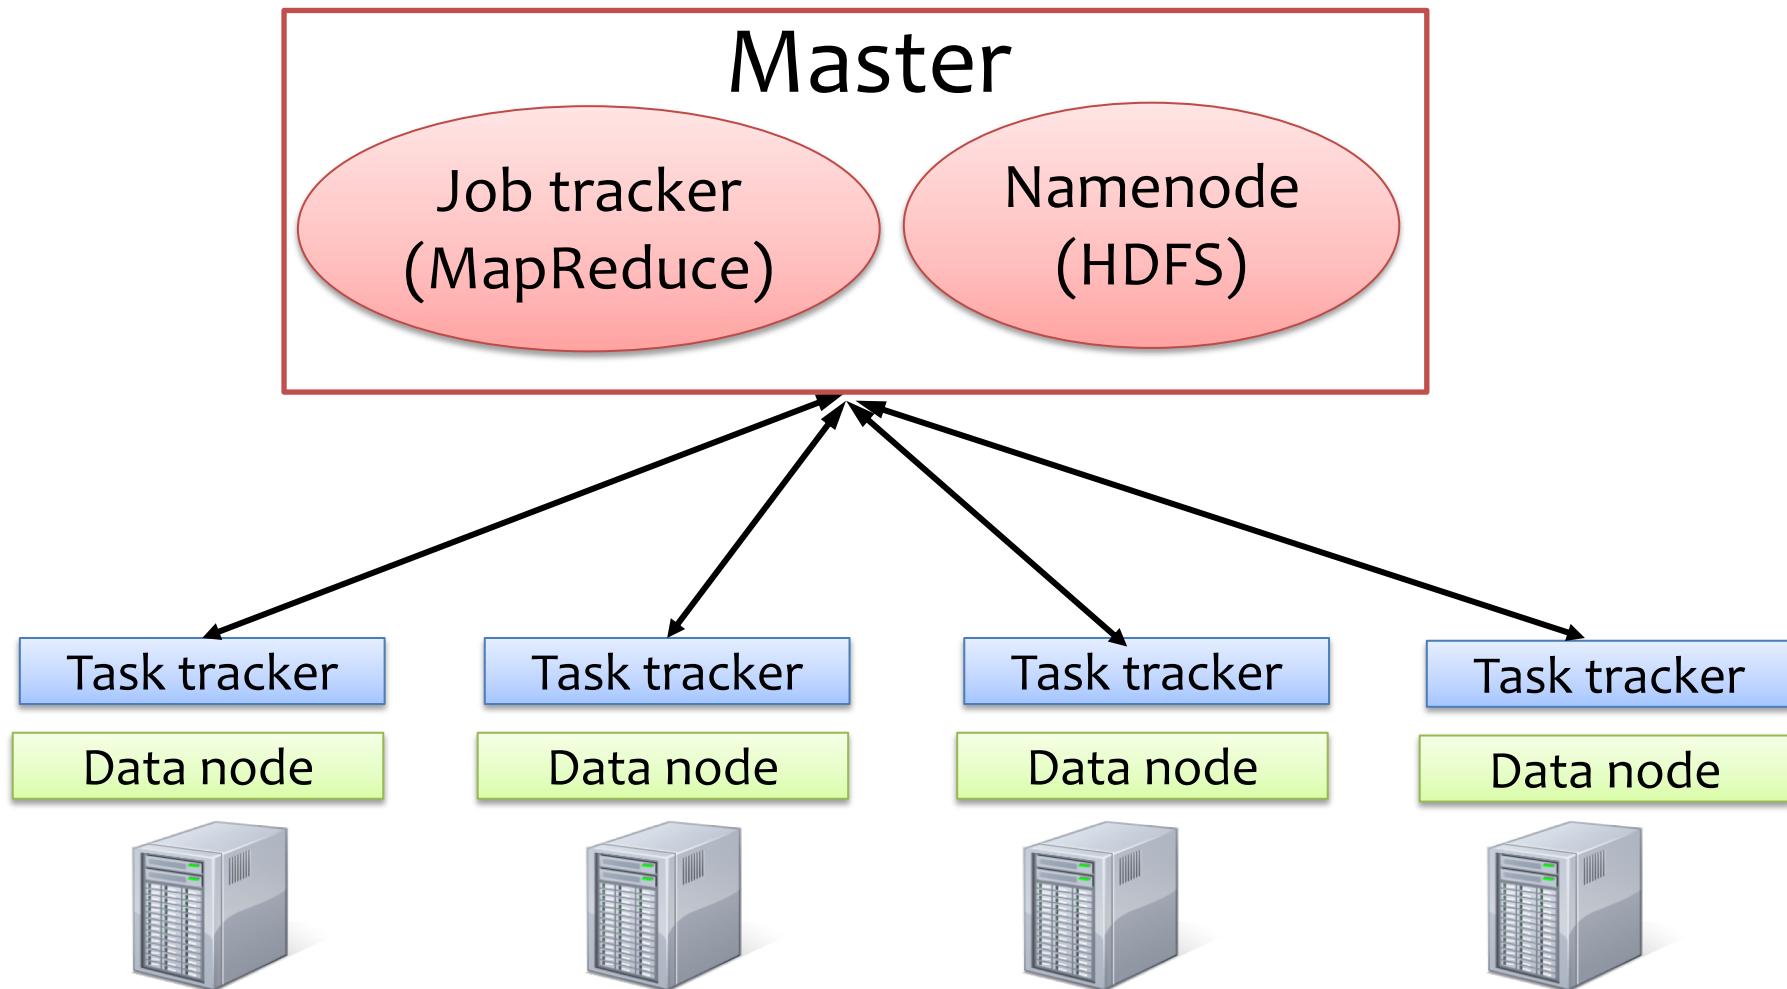
Word-count example



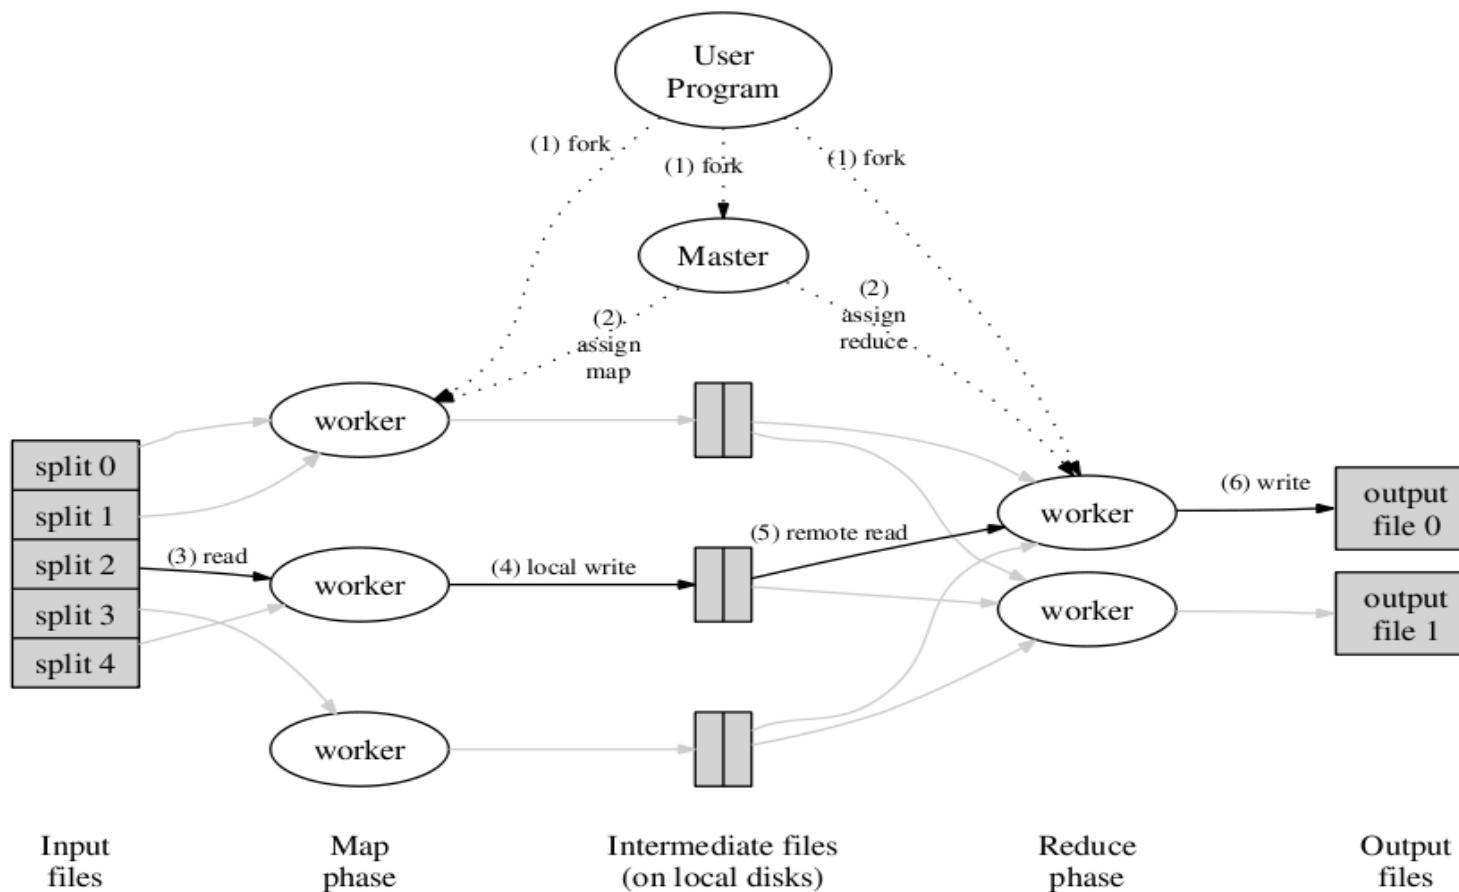
MapReduce software stack



MapReduce software stack



Runtime execution



Design challenges revisited

- Parallelization
- Communication
- Synchronization
- Load balancing
- Faults & semantics
- Scheduling

References

- Compulsory reading:
 - MapReduce [OSDI'04]:
<https://static.googleusercontent.com/media/research.google.com/en//archive/mapreduce-osdi04.pdf>
- Recommended reading:
 - YARN [SoCC'13]
 - the next generation of M/R
 - Dryad [EuroSys'07]
 - Generalized framework for data-parallel computations
 - Spark [NSDI'12]
 - In-memory distributed data parallel computing