

Massively Available Key-Value Stores & Consistent Hashing

Key-value stores

- *What is a key-value-store?*
- *Why are key-value stores needed?*
- Key-value-store client interface
- Key-value stores in practice
- Common features & non-features

What mechanisms make them work?

What are key-value stores?

- Container for key-value pairs (databases)
- Distributed, multi-component, systems
- NoSQL semantics (non-relational)
- KV-stores offer **simpler query semantics** in exchange for **increased scalability, speed, availability, and flexibility**
- Data model not new



DBMS (SQL)

Students Table

Student	ID*
John Smith	084
Jane Bloggs	100
John Smith	182
Mark Antony	219

Activities Table

ID*	Activity*	Cost
084	Swimming	\$17
084	Tennis	\$36
100	Squash	\$40
100	Swimming	\$17
182	Tennis	\$36
219	Golf	\$47
219	Swimming	\$15
219	Squash	\$40

- Relational data schema
- Data types
- Foreign keys
- Full SQL support

Key-value store

Key	Value
John Smith	{Activity:Name=Swimming}
Jane Bloggs	{Activity:Cost=57}
Mark Anthony	{ID=219}

- No data schema
- Raw byte access
- No relations
- Single-row operations

Why are key-value stores needed?

- Today's internet applications
 - Huge amounts of stored data
 - Huge number of Internet users
 - Frequent updates
 - Fast retrieval of information
 - Rapidly changing data definitions
- Ever more users, ever more data



Why are key-value stores needed?

- Horizontal scalability
 - User growth, traffic patterns change
 - Adapt to number of requests, data size
- Performance
 - High speed for single-record read and write operations
- Flexibility
 - Adapt to changing data definitions

Why are key-value stores needed?

- Reliability
 - Thousands of components at play
 - Uses commodity hardware: failure is the norm
 - Provide failure recovery
- Availability and geo-distribution
 - Users are worldwide
 - Guarantee fast access

Key-value store client interface

- Main operations
 - Write/update **put**(key, value)
 - Read **get**(key)
 - Delete **delete**(key)
- Usually no aggregation, no table joins, no transactions!

Hbase: Key-value store client interface

```
Configuration conf = HBaseConfiguration.create();
conf.set("hbase.zookeeper.quorum", "192.168.127.129");

HTable table = new HTable(conf, „MyBaseTable");

Put put = new Put(Bytes.toBytes("key1"));
put.add(Bytes.toBytes("colfam1"), Bytes.toBytes(„value"), Bytes.toBytes(200));
table.put(put);

Get get = new Get(Bytes.toBytes("key1"));
Result result = table.get(get);
byte[] val = result.getValue(Bytes.toBytes("colfam1"), Bytes.toBytes(„value"));
System.out.println("Value: " + Bytes.toInt(val));
```

Hbase: Key-value store client interface

```
Configuration conf = HBaseConfiguration.create();  
conf.set("hbase.zookeeper.quorum", "192.168.127.129");
```

```
HTable table = new HTable(conf, „MyBaseTable“);
```

```
Put put = new Put(Bytes.toBytes("key1"));  
put.add(Bytes.toBytes("colfam1"), Bytes.toBytes(„value“), Bytes.toBytes(200));  
table.put(put);
```

```
Get get = new Get(Bytes.toBytes("key1"));  
Result result = table.get(get);  
byte[] val = result.getValue(Bytes.toBytes("colfam1"), Bytes.toBytes(„value“));  
System.out.println("Value: " + Bytes.toInt(val));
```



Initialization
Using
ZooKeeper

Hbase: Key-value store client interface

```
Configuration conf = HBaseConfiguration.create();  
conf.set("hbase.zookeeper.quorum", "192.168.127.129");
```

Initialization
Using
ZooKeeper

```
HTable table = new HTable(conf, „MyBaseTable“);
```

Column Family:
“Schema”

```
Put put = new Put(Bytes.toBytes("key1"),  
put.add(Bytes.toBytes("colfam1"), Bytes.toBytes("value"), Bytes.toBytes(200));  
table.put(put);
```

```
Get get = new Get(Bytes.toBytes("key1"));  
Result result = table.get(get);  
byte[] val = result.getValue(Bytes.toBytes("colfam1"), Bytes.toBytes("value"));  
System.out.println("Value: " + Bytes.toInt(val));
```

Hbase: Key-value store client interface

```
Configuration conf = HBaseConfiguration.create();  
conf.set("hbase.zookeeper.quorum", "192.168.127.129");
```

Initialization
Using
ZooKeeper

```
HTable table = new HTable(conf, „MyBaseTable“);
```

Column Family:
“Schema”

```
Put put = new Put(Bytes.toBytes("key1"),  
put.add(Bytes.toBytes("colfam1"), Bytes.toBytes("value"), Bytes.toBytes(200));  
table.put(put);
```

Column:
Defined at run-time
(“wide column”
stores)

```
Get get = new Get(Bytes.toBytes("key1"));  
Result result = table.get(get);  
byte[] val = result.getValue(Bytes.toBytes("colfam1"), Bytes.toBytes("value"));  
System.out.println("Value: " + Bytes.toInt(val));
```

Key-value store in practice

- BigTable
- Apache HBase
- Apache Cassandra
- Redis
- Amazon Dynamo
- Yahoo! PNUTS

Google

facebook




amazon

YAHOO!

Common elements of key-value stores

- Failure detection, failure recovery
- Replication (*cf. Replication*)
 - Store and manage multiple copies of data
- Versioning (*cf. Time*)
 - Store different versions of data
 - Timestamping


Consistent Hashing



Organization-Based Analysis of Web-Object Sharing and Caching
Alec Wolma et al.


Caching

- Five cache nodes
- Client determines node to contact




Caching

- Five cache nodes
- Client determines node to contact




Caching

- Five cache nodes
- Client determines node to contact




Caching

- Five cache nodes
- Client determines node to contact



Caching

- Five cache nodes
- Client determines node to contact



Problem: Mapping objects to caches

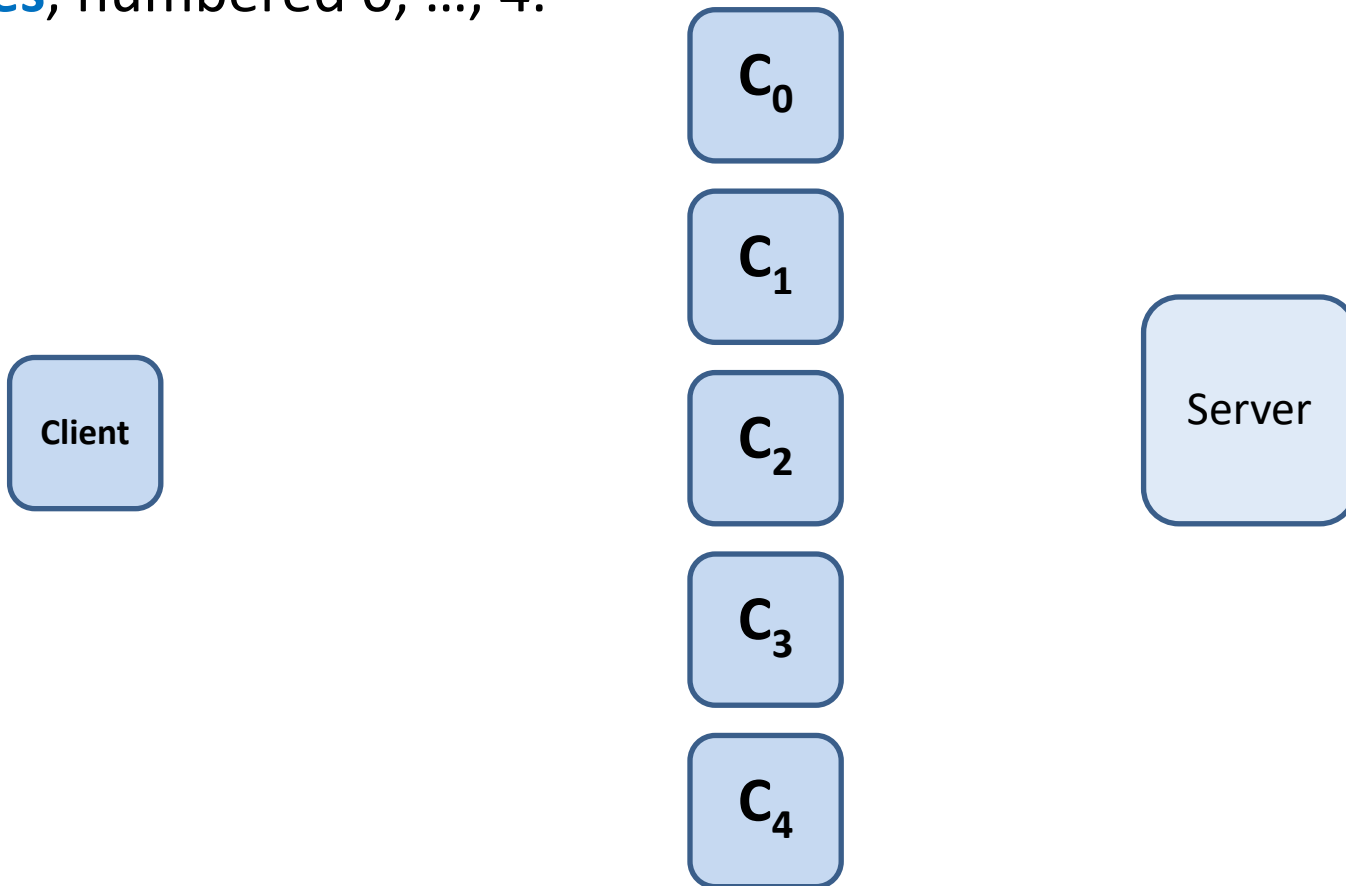
- Given a number of caches (e.g., cooperative caching, CDNs, etc.)
- Each cache should carry an **equal share of objects**
- **Clients** need to **know what cache to query** for a given object
- **Horizontally partition** (shard) object ID space
 - **Doesn't work** with **skewed distributions**: e.g., 10 servers, each handles 100 IDs, but all objects have IDs between 1-100 or 900-1000
- **Caches** should be able to **come and go** without disrupting the whole operation (i.e., non-effected caches)

Solution attempt: Use hashing

- **Map object ID** (e.g., URL u) **into one of the caches**
- Use a hash function that **maps u to node $h(u)$**
 - For example, $h(x) = (ax + b) \bmod p$, where p is range of $h(x)$, i.e., the number of caches
 - Interpret u as a number based on bit pattern of object ID (or URL)
- Hashing tends to **distribute input uniformly** across range of hash function
 - Objects (URLs) are **equally balanced** across caches, even if object IDs are skewed (i.e., highly clustered in ID space)
- No one cache responsible for an **uneven share of objects/URLs**
- No disproportionately loaded node (potential bottleneck)


$$h(u) = (7u + 4) \bmod 5$$

Assume, we have **five**
caches, numbered 0, ..., 4.




$$h(u) = (7u + 4) \bmod 5$$

Assume, we have **five**
caches, numbered 0, ..., 4.




$$h(u) = (7u + 4) \bmod 5$$

Assume, we have **five**
caches, numbered 0, ..., 4.



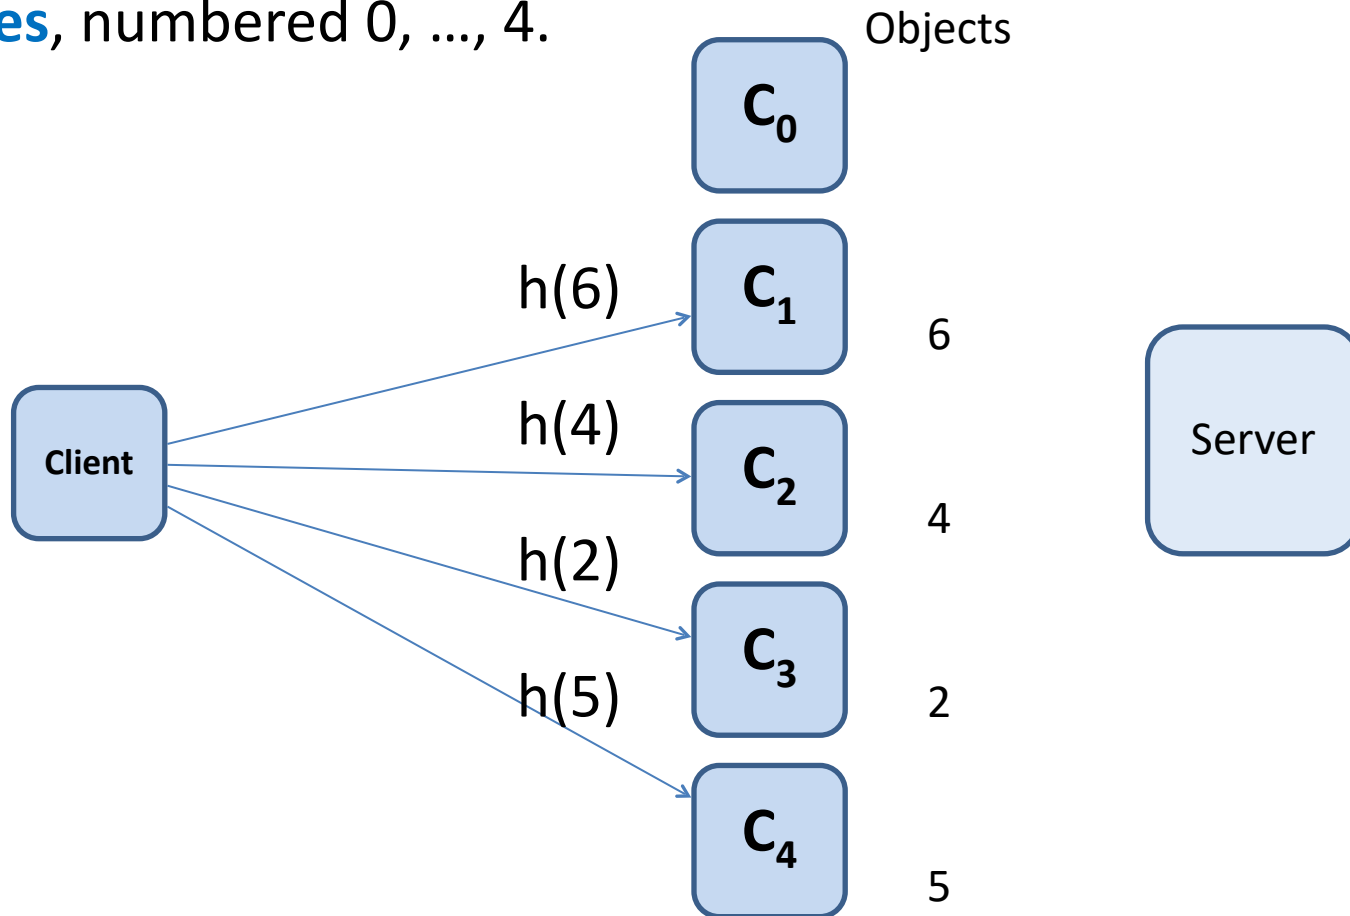
$$h(u) = (7u + 4) \bmod 5$$

Assume, we have **five**
caches, numbered 0, ..., 4.




$$h(u) = (7u + 4) \bmod 5$$

Assume, we have **five caches**, numbered 0, ..., 4.



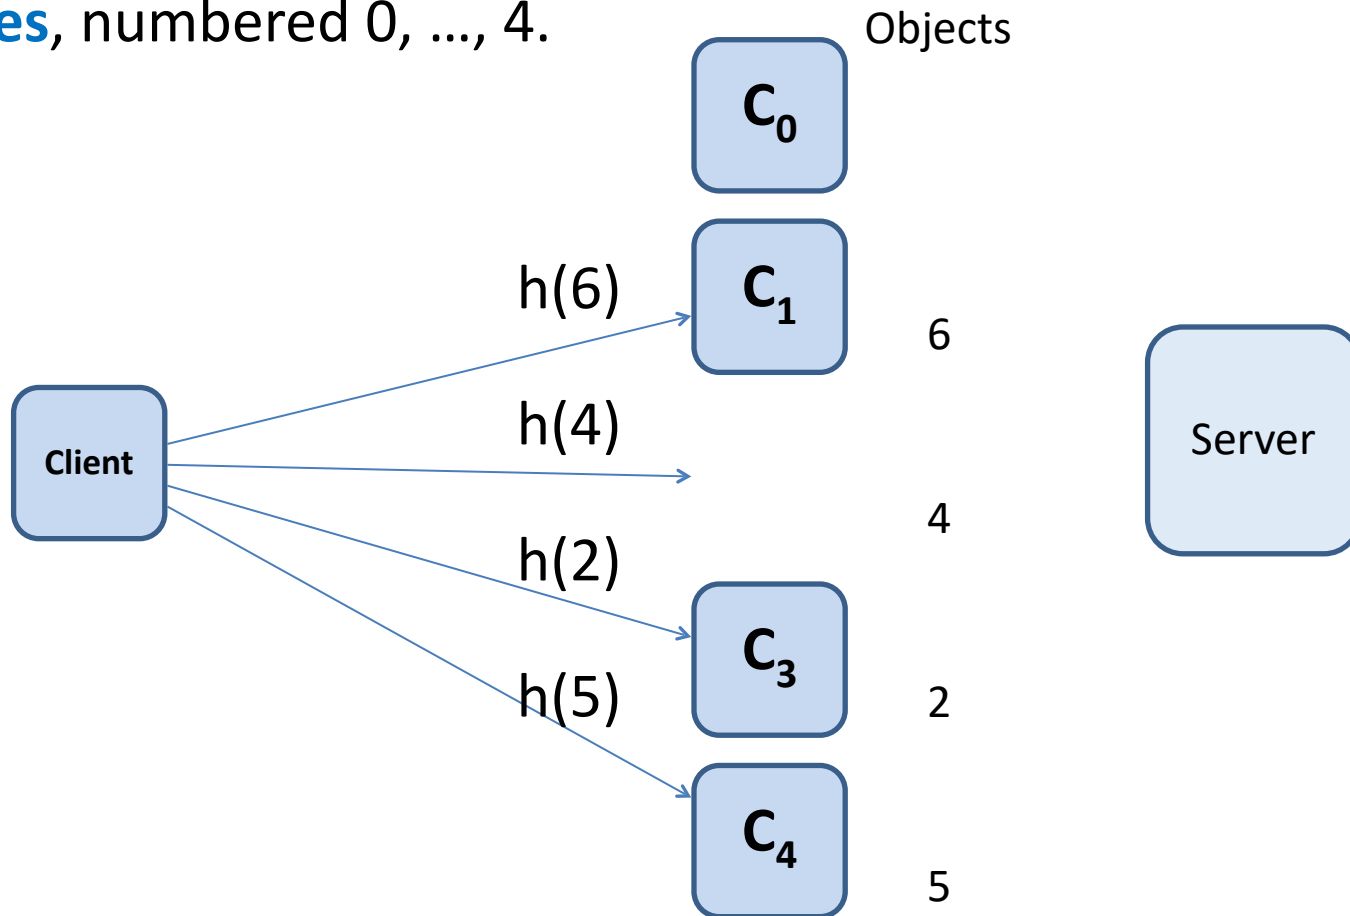
$$h(u) = (7u + 4) \bmod 5$$

Assume, we have **five**
caches, numbered 0, ..., 4.




$$h(u) = (7u + 4) \bmod 5$$

Assume, we have **five caches**, numbered 0, ..., 4.




$h(u) = (7u + 4) \bmod 4$
 (now have to map across 4 caches)



$h(u) = (7u + 4) \bmod 4$
(now have to map across 4 caches)


**Removing a cache
changes location
of almost every objects!**



$$h(u) = 7u + 4 \bmod 4$$


(mapped across 4 nodes)

Objects



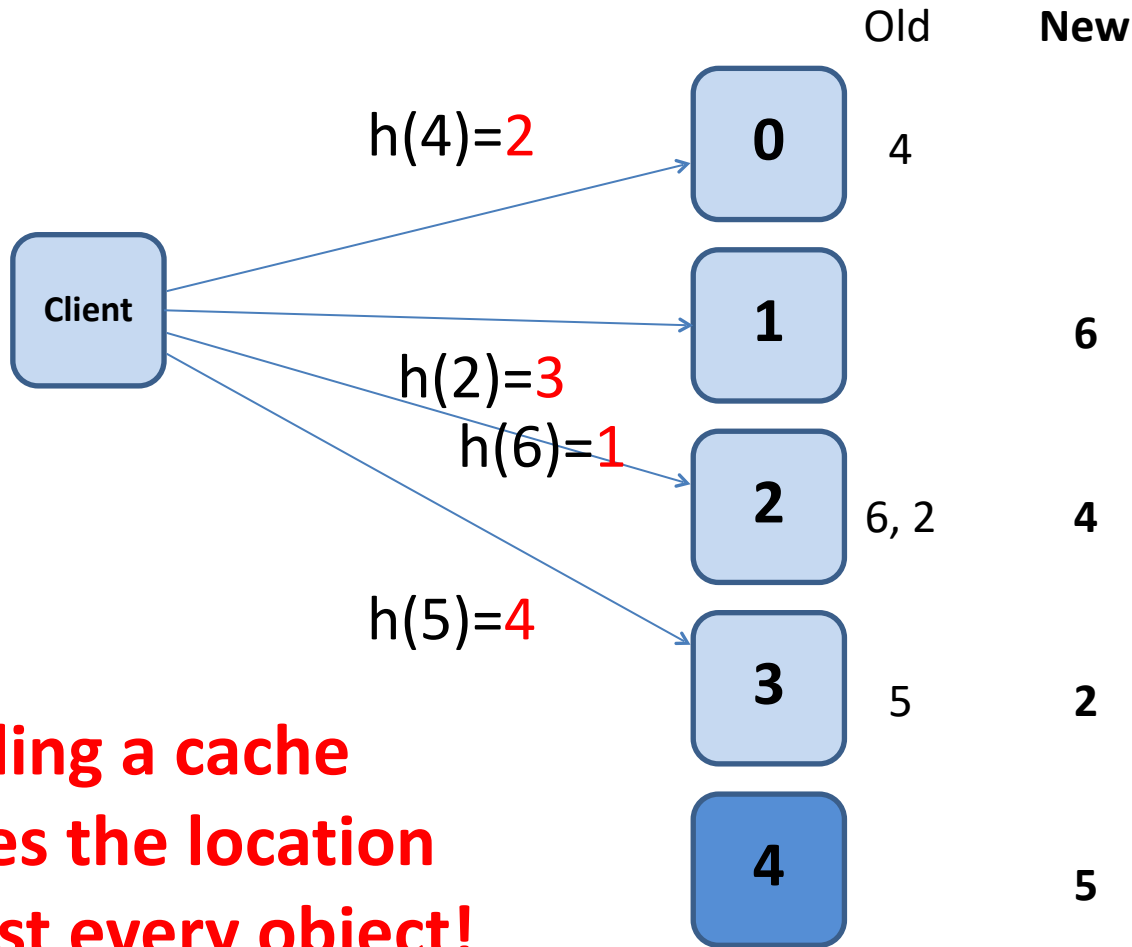
$$h(u) = (7u + 4) \bmod 5$$

(adding a cache again)



$$h(u) = (7u + 4) \bmod 5$$

(adding a cache again)



**Adding a cache
changes the location
of almost every object!**


Consistent hashing

- **Goals**
 - Uniform distribution of objects across nodes
 - Easily find objects
 - Let any client perform a local computation mapping a URL to node that contains referenced object
 - **Allow for nodes to be added/removed without much disruption** – remap only n/m objects (n objects, m slots)
- D. Karger *et al.*, MIT, 1997
- Basis for Akamai
 - CDN company (content delivery network)
 - Web cache as a service

Consistent hashing

Key idea intuition


- Select a **base hash function** that maps input identifier to the number range $[0, \dots, m-1]$
- *E.g.*, $h(x) = (ax + b) \bmod m$
- Interpret range of $h(..)$ as array that wraps around (i.e., a circle)
- $h(..)$ gives slot in array (circle) and wraps around at $m-1$ to 0
- Each **object** is mapped to a **slot** via $h(..)$
- Each **cache** is mapped to a **slot** via $h(..)$
- Assign each object to the closest cache slot in clockwise direction on the circle



Consistent hashing

Key idea intuition


- Select a **base hash function** that maps input identifier to the number range $[0, \dots, m-1]$
- *E.g.*, $h(x) = (ax + b) \bmod m$
- Interpret range of $h(..)$ as array that wraps around (i.e., a circle)
- $h(..)$ gives slot in array (circle) and wraps around at $m-1$ to 0
- Each **object** is mapped to a **slot** via $h(..)$
- Each **cache** is mapped to a **slot** via $h(..)$
- Assign each object to the closest **cache slot** in **clockwise direction** on the circle



Consistent hashing

Key idea intuition


- Select a **base hash function** that maps input identifier to the number range $[0, \dots, m-1]$
- *E.g.*, $h(x) = (ax + b) \bmod m$
- Interpret range of $h(..)$ as array that wraps around (i.e., a circle)
- $h(..)$ gives slot in array (circle) and wraps around at $m-1$ to 0
- Each **object** is mapped to a **slot** via $h(..)$
- Each **cache** is mapped to a **slot** via $h(..)$
- Assign each object to the closest **cache slot** in **clockwise direction** on the circle



Consistent hashing

Key idea intuition


- Select a **base hash function** that maps input identifier to the number range $[0, \dots, m-1]$
- *E.g.*, $h(x) = (ax + b) \bmod m$
- Interpret range of $h(..)$ as array that wraps around (i.e., a circle)
- $h(..)$ gives slot in array (circle) and wraps around at $m-1$ to 0
- Each **object** is mapped to a **slot** via $h(..)$
- Each **cache** is mapped to a **slot** via $h(..)$
- Assign each object to the closest **cache slot** in **clockwise direction** on the circle



Consistent hashing

Original interpretation


- Select a **base hash function** that maps input identifier to the number range $[0, \dots, M]$
- Divide by M , re-mapping $[0, \dots, M]$ to $[0, 1]$
- Interpret this interval as the **unit circle**: Here, circle with circumference 1 (normally radius 1)
- Each **object** is mapped to a **point** on unit circle via $h(..)$
- **Each cache** is mapped to a **point** on unit circle via $h(..)$
- Assign **each URL** to **closest cache point in clockwise direction** on the circle



Consistent hashing

Original interpretation


- Select a **base hash function** that maps input identifier to the number range $[0, \dots, M]$
- Divide by M , re-mapping $[0, \dots, M]$ to $[0, 1]$
- Interpret this interval as the **unit circle**: Here, circle with circumference 1 (normally radius 1)
- Each **object** is mapped to a **point** on unit circle via $h(..)$
- **Each cache is mapped to a point** on unit circle via $h(..)$
- Assign **each URL to closest cache point in clockwise direction** on the circle



Consistent hashing

Original interpretation


- Select a **base hash function** that maps input identifier to the number range $[0, \dots, M]$
- Divide by M , re-mapping $[0, \dots, M]$ to $[0, 1]$
- Interpret this interval as the **unit circle**: Here, circle with circumference 1 (normally radius 1)
- Each **object** is mapped to a **point** on unit circle via $h(..)$
- **Each cache** is mapped to a **point** on unit circle via $h(..)$
- Assign **each URL** to **closest cache point in clockwise direction** on the circle



Consistent hashing

Original interpretation


- Select a **base hash function** that maps input identifier to the number range $[0, \dots, M]$
- Divide by M , re-mapping $[0, \dots, M]$ to $[0, 1]$
- Interpret this interval as the **unit circle**: Here, circle with circumference 1 (normally radius 1)
- Each **object** is mapped to a **point** on unit circle via $h(..)$
- **Each cache** is mapped to a **point** on unit circle via $h(..)$
- Assign **each URL** to **closest cache point in clockwise direction** on the circle



Consistent hashing

Original interpretation


- Select a **base hash function** that maps input identifier to the number range $[0, \dots, M]$
- Divide by M , re-mapping $[0, \dots, M]$ to $[0, 1]$
- Interpret this interval as the **unit circle**: Here, circle with circumference 1 (normally radius 1)
- Each **object** is mapped to a **point** on unit circle via $h(..)$
- **Each cache** is mapped to a **point** on unit circle via $h(..)$
- Assign **each URL** to **closest cache point in clockwise direction** on the circle



Consistent hashing

Original interpretation


- Select a **base hash function** that maps input identifier to the number range $[0, \dots, M]$
- Divide by M , re-mapping $[0, \dots, M]$ to $[0, 1]$
- Interpret this interval as the **unit circle**: Here, circle with circumference 1 (normally radius 1)
- Each **object** is mapped to a **point** on unit circle via $h(..)$
- **Each cache** is mapped to a **point** on unit circle via $h(..)$
- Assign **each URL** to **closest cache point in clockwise direction** on the circle



Consistent hashing

Original interpretation


- Select a **base hash function** that maps input identifier to the number range $[0, \dots, M]$
- Divide by M , re-mapping $[0, \dots, M]$ to $[0, 1]$
- Interpret this interval as the **unit circle**: Here, circle with circumference 1 (normally radius 1)
- Each **object** is mapped to a **point** on unit circle via $h(..)$
- **Each cache is mapped to a point** on unit circle via $h(..)$
- Assign **each URL to closest cache point in clockwise direction** on the circle



Consistent hashing

Original interpretation


- Select a **base hash function** that maps input identifier to the number range $[0, \dots, M]$
- Divide by M , re-mapping $[0, \dots, M]$ to $[0, 1]$
- Interpret this interval as the **unit circle**: Here, circle with circumference 1 (normally radius 1)
- Each **object** is mapped to a **point** on unit circle via $h(..)$
- **Each cache** is mapped to a **point** on unit circle via $h(..)$
- Assign **each URL** to **closest cache point in clockwise direction** on the circle



Consistent hashing

Original interpretation


- Select a **base hash function** that maps input identifier to the number range $[0, \dots, M]$
- Divide by M , re-mapping $[0, \dots, M]$ to $[0, 1]$
- Interpret this interval as the **unit circle**: Here, circle with circumference 1 (normally radius 1)
- Each **object** is mapped to a **point** on unit circle via $h(..)$
- **Each cache** is mapped to a **point** on unit circle via $h(..)$
- Assign **each URL** to **closest cache point in clockwise direction** on the circle




Consistent hashing

Original interpretation

- Select a **base hash function** that maps input identifier to the number range $[0, \dots, M]$
- Divide by M , re-mapping $[0, \dots, M]$ to $[0, 1]$
- Interpret this interval as the **unit circle**: Here, circle with circumference 1 (normally radius 1)
- Each **object** is mapped to a **point** on unit circle via $h(..)$
- **Each cache** is mapped to a **point** on unit circle via $h(..)$
- Assign **each URL** to **closest cache point in clockwise direction** on the circle




Mapping items to caches



Items 2, 3 mapped to **B**
Items 4, 5, 6 mapped to **C**
Items 7, 1 mapped to **A**


● Cache
● Object

Removing a cache




Items	2, 3	mapped to B
Items	4, 5, 6	mapped to C
Items	7, 1	mapped to A

Removing a cache




Items	2, 3	mapped to B
Items	4, 5, 6	mapped to C
Items	7, 1	mapped to A

Removing a cache



Items 2, 3, **7**, **1** mapped to **B**
Items 4, 5, 6 mapped to **C**


Adding a cache



Items 7, 1, 2, 3 mapped to B


Items 4, 5, 6 mapped to **C**

Adding a cache




Items 7, 1, 2, 3 mapped to B
Items 4, 5, 6 mapped to C

Adding a cache



Items	3	mapped to B
Items	4, 5, 6	mapped to C
Items	7, 1, 2	mapped to A


Processing a Lookup (key)



Information about
node addition & removal
(e.g., via gossiping or via a
coordination service)



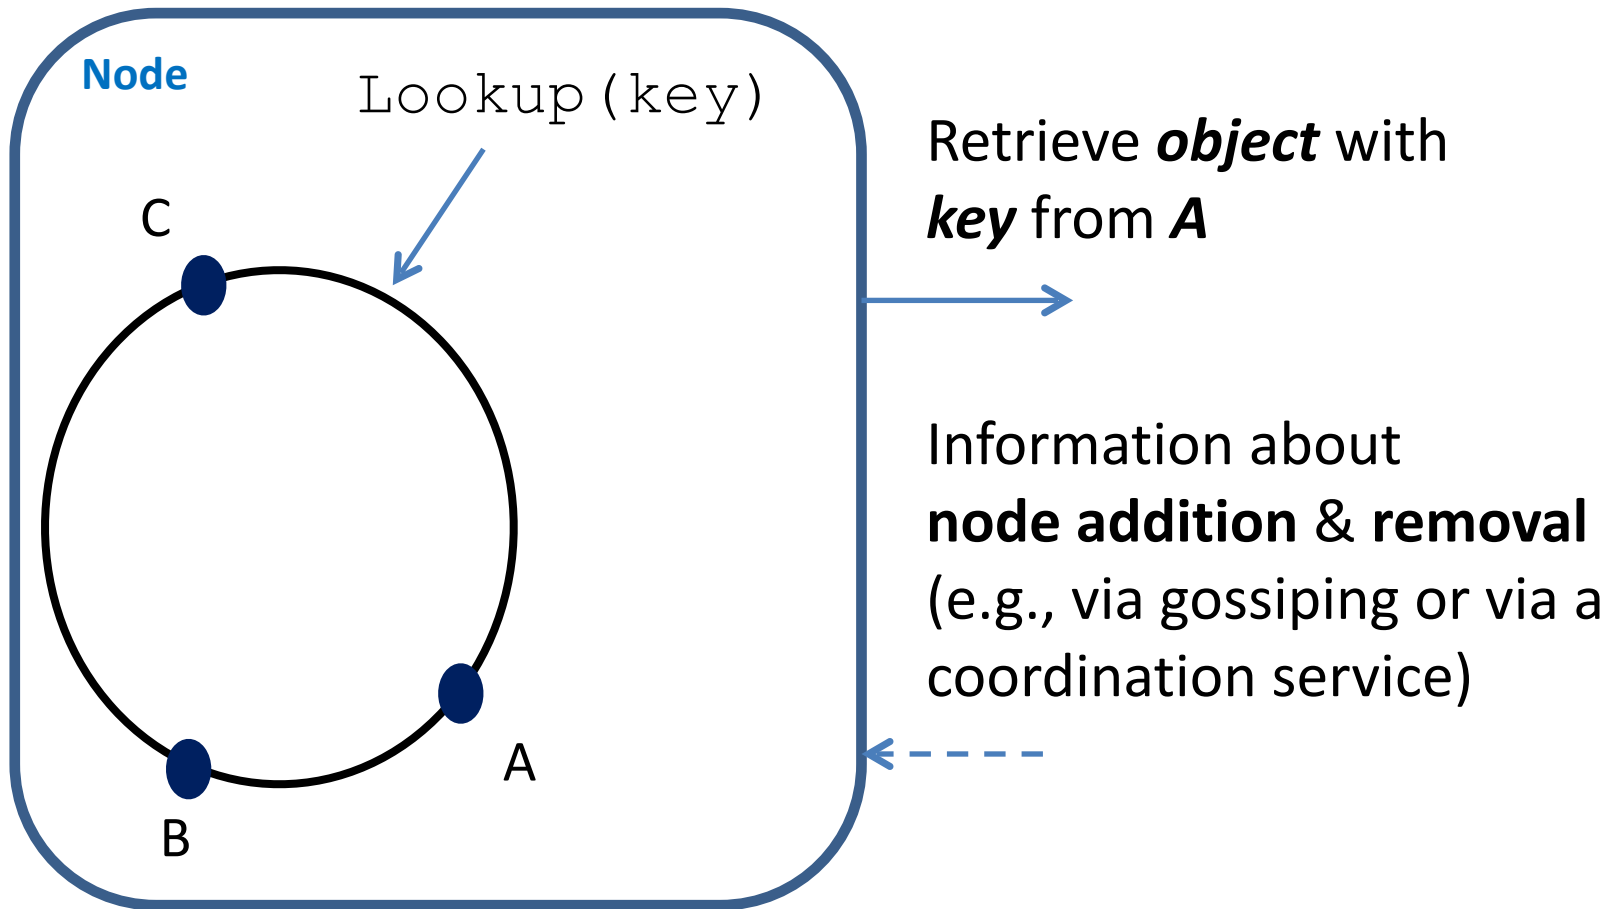
Processing a Lookup (key)



Information about
node addition & removal
(e.g., via gossiping or via a
coordination service)




Processing a Lookup (key)



Cache lookup data structure at each node

- Store **cache points** in a **binary tree**
- Find **clockwise successor of a URL point** by single search in tree (takes **$O(\log n)$ time**)
- For a constant time technique, cf. Karger et al., 1997

$\text{key}[\text{left}(x)] \leq \text{key}[x] \leq \text{key}[\text{right}(x)]$



Base hash function: MD5

- **Message Digest 5** (MD5), R. Rivest, **1992** (MD1, ..., MD6)
- **Hash function** that produces a **128-bit** (16-byte) **hash value**
- Maps variable-length message into a **fixed-length output**
- MD5 hash is typically expressed as a hex number (32 digits)
- It's been shown that **MD5 is not collision resistant**
- US-CERT about MD5 “*should be considered cryptographically broken and unsuitable for further use*” (for security, not for caching)
- SHA-2 is a more appropriate cryptographic hash function
- For consistent hashing, MD5 is sufficient

MD5 examples

- MD5(*"The quick brown fox jumps over the lazy dog"*) = 9e107d9d372bb6826bd81d3542a419d6
- MD5(*"The quick brown fox jumps over the lazy dog."*) = e4d909c290d0fb1ca068ffaddf22cbd0
- MD5(*"*) = d41d8cd98f00b204e9800998ecf8427e

MD5 examples

- MD5(*"The quick brown fox jumps over the lazy dog"*) = 9e107d9d372bb6826bd81d3542a419d6
- MD5(*"The quick brown fox jumps over the lazy dog."*) = e4d909c290d0fb1ca068ffaddf22cbd0
- MD5(*"*) = d41d8cd98f00b204e9800998ecf8427e

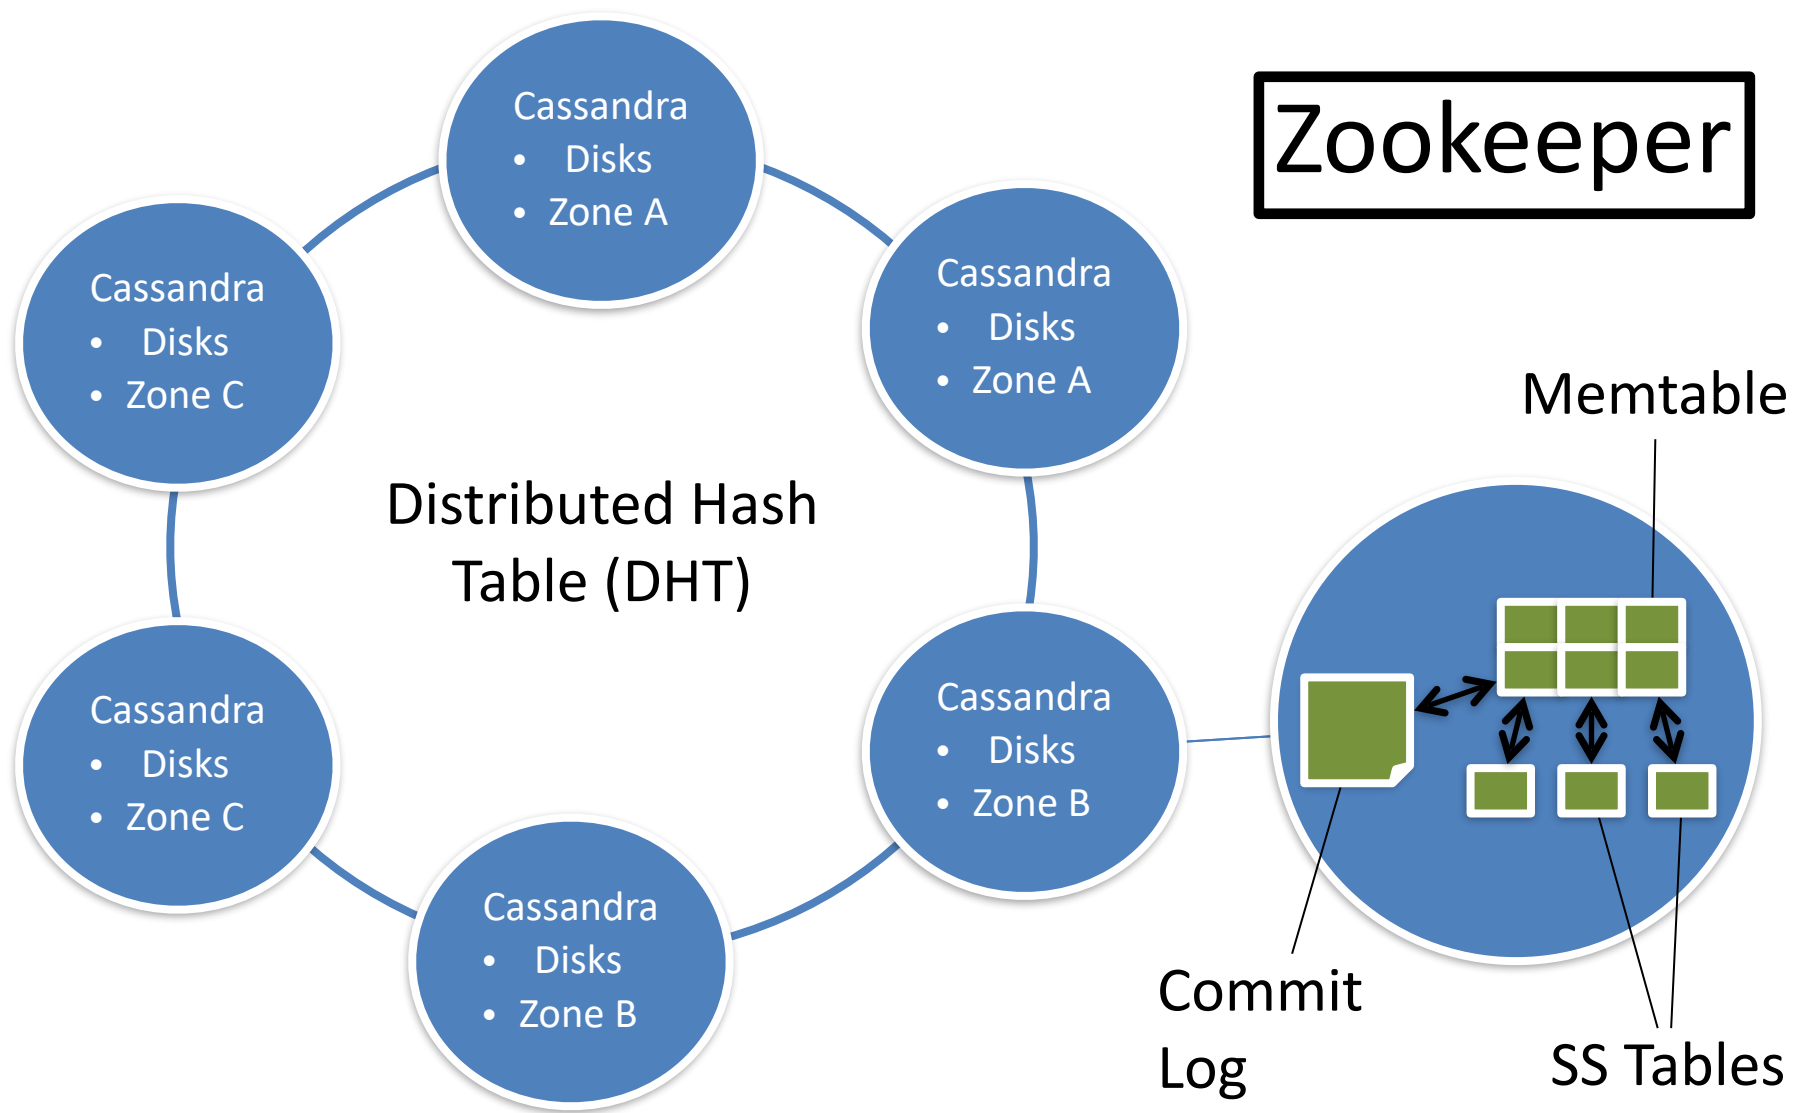


DYNAMO / CASSANDRA

Cassandra

- Developed by Facebook
- Based on Amazon Dynamo (but open-source)
- Structured storage nodes (**no GFS** used)
- **Decentralized** architecture (no master assignment)
- **Consistent hashing** for load balancing
- Eventual consistency
- **Gossiping** to exchange information

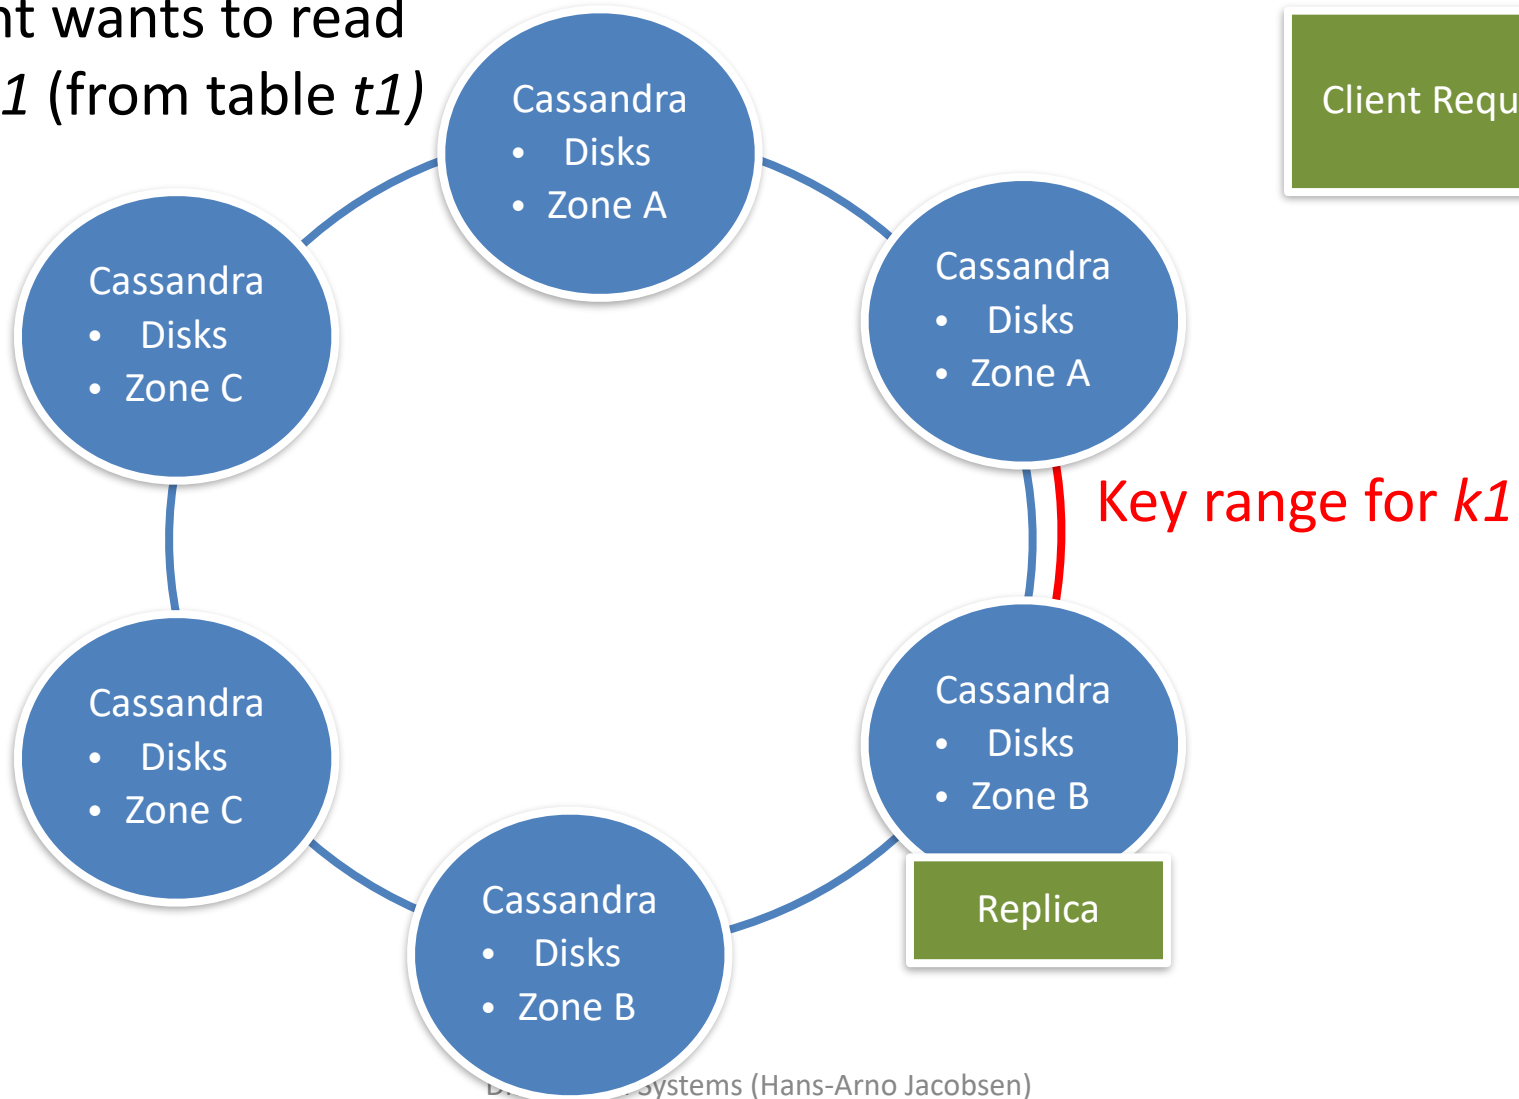
Cassandra architecture overview



Cassandra global read-path

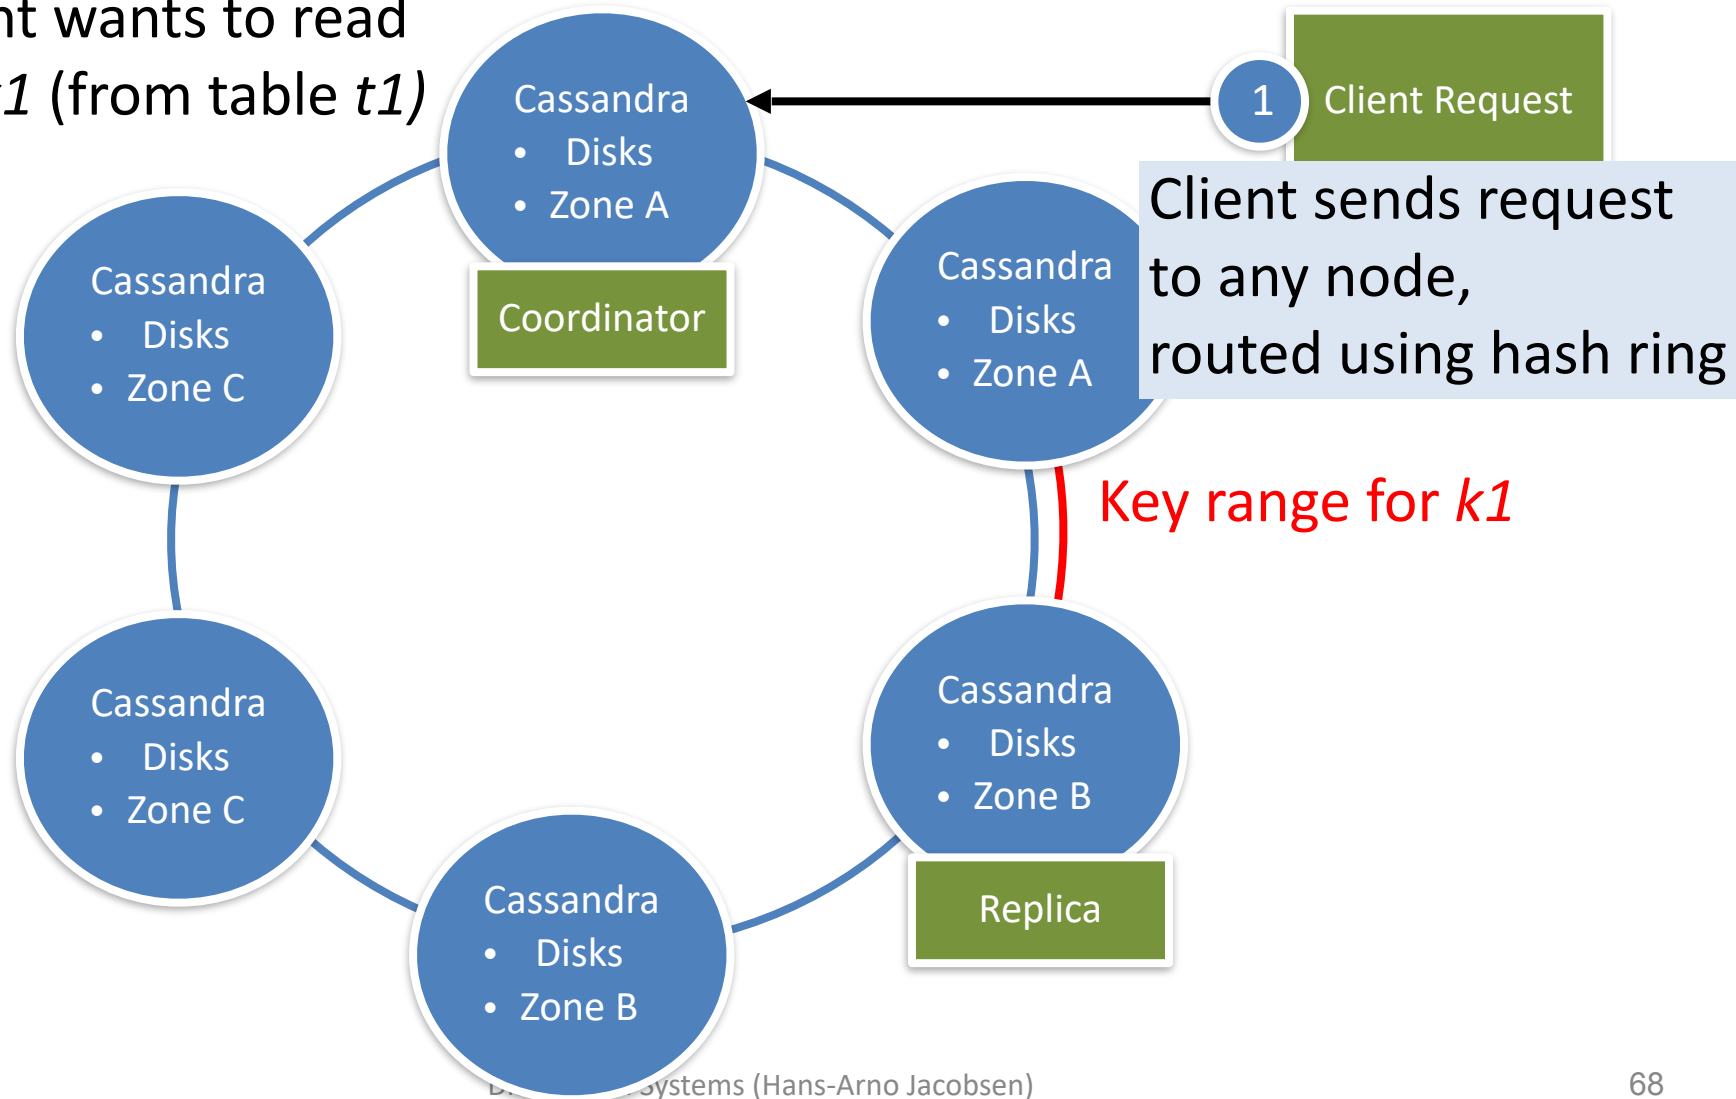
Client wants to read
key $k1$ (from table $t1$)

Client Request



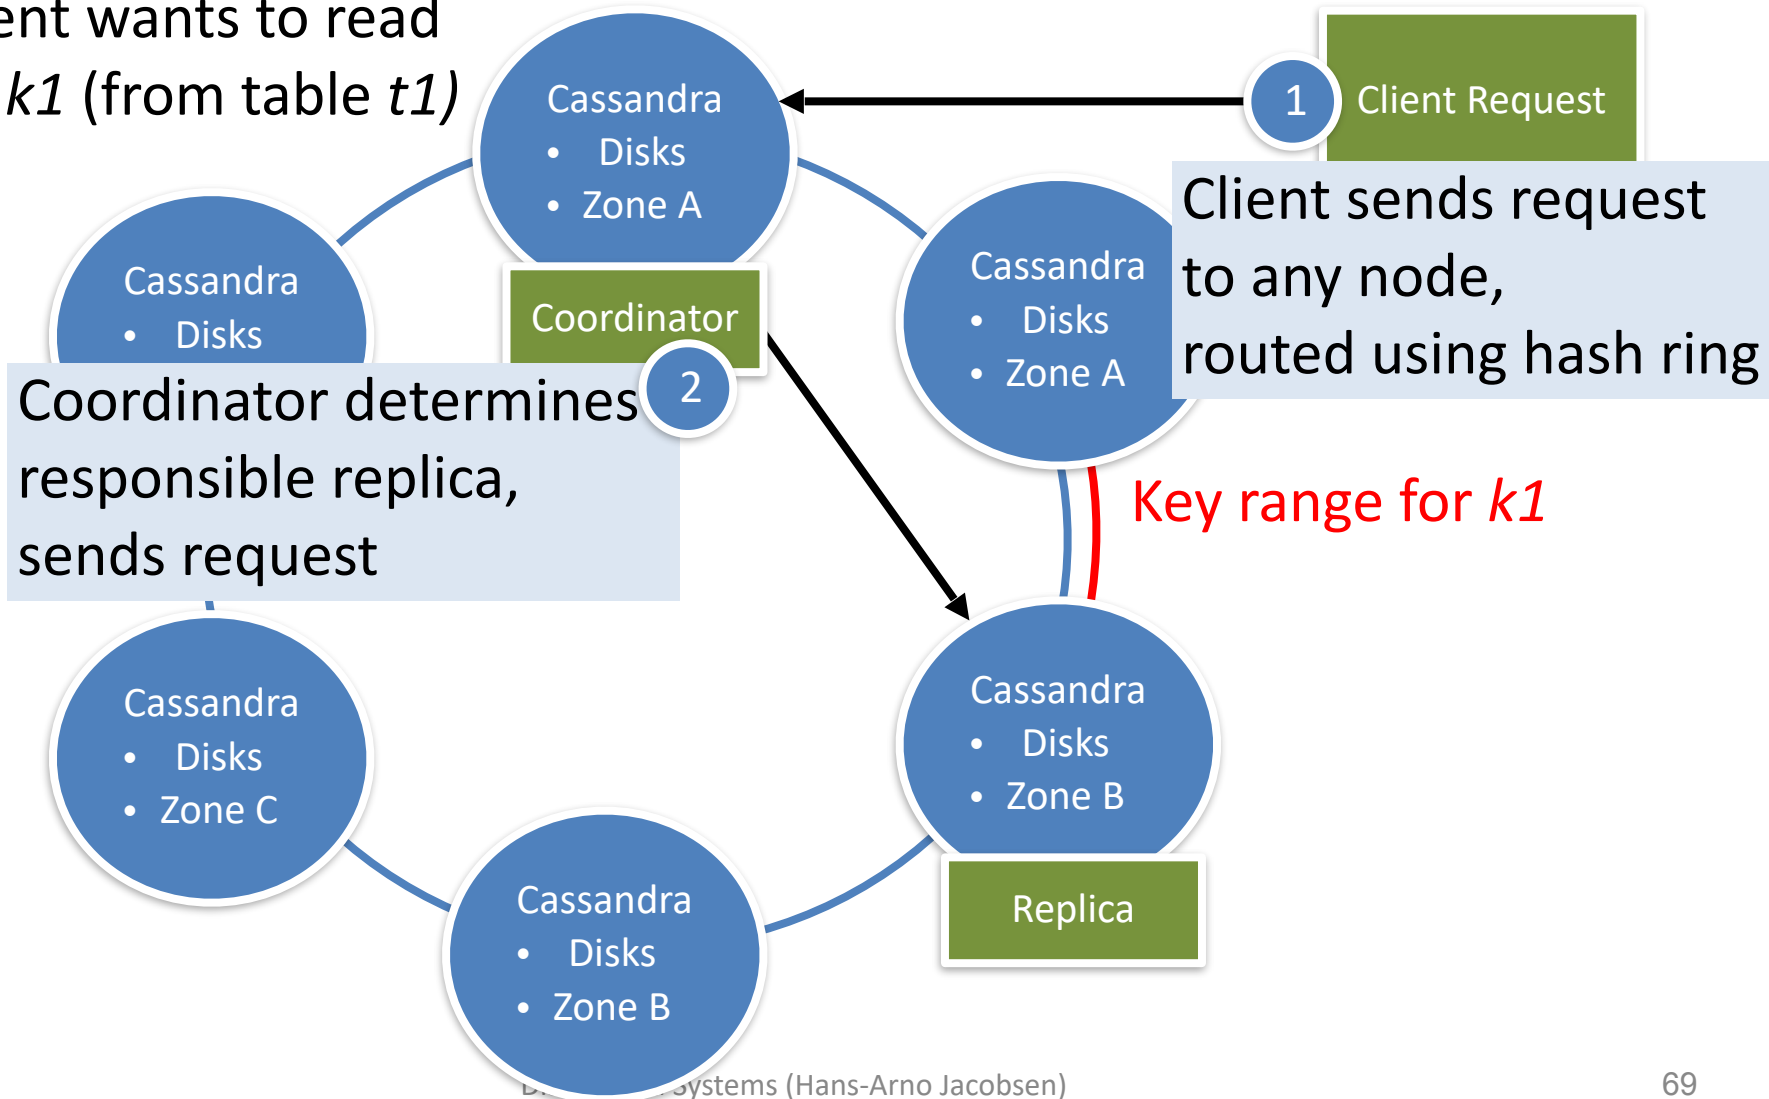
Cassandra global read-path

Client wants to read
key $k1$ (from table $t1$)



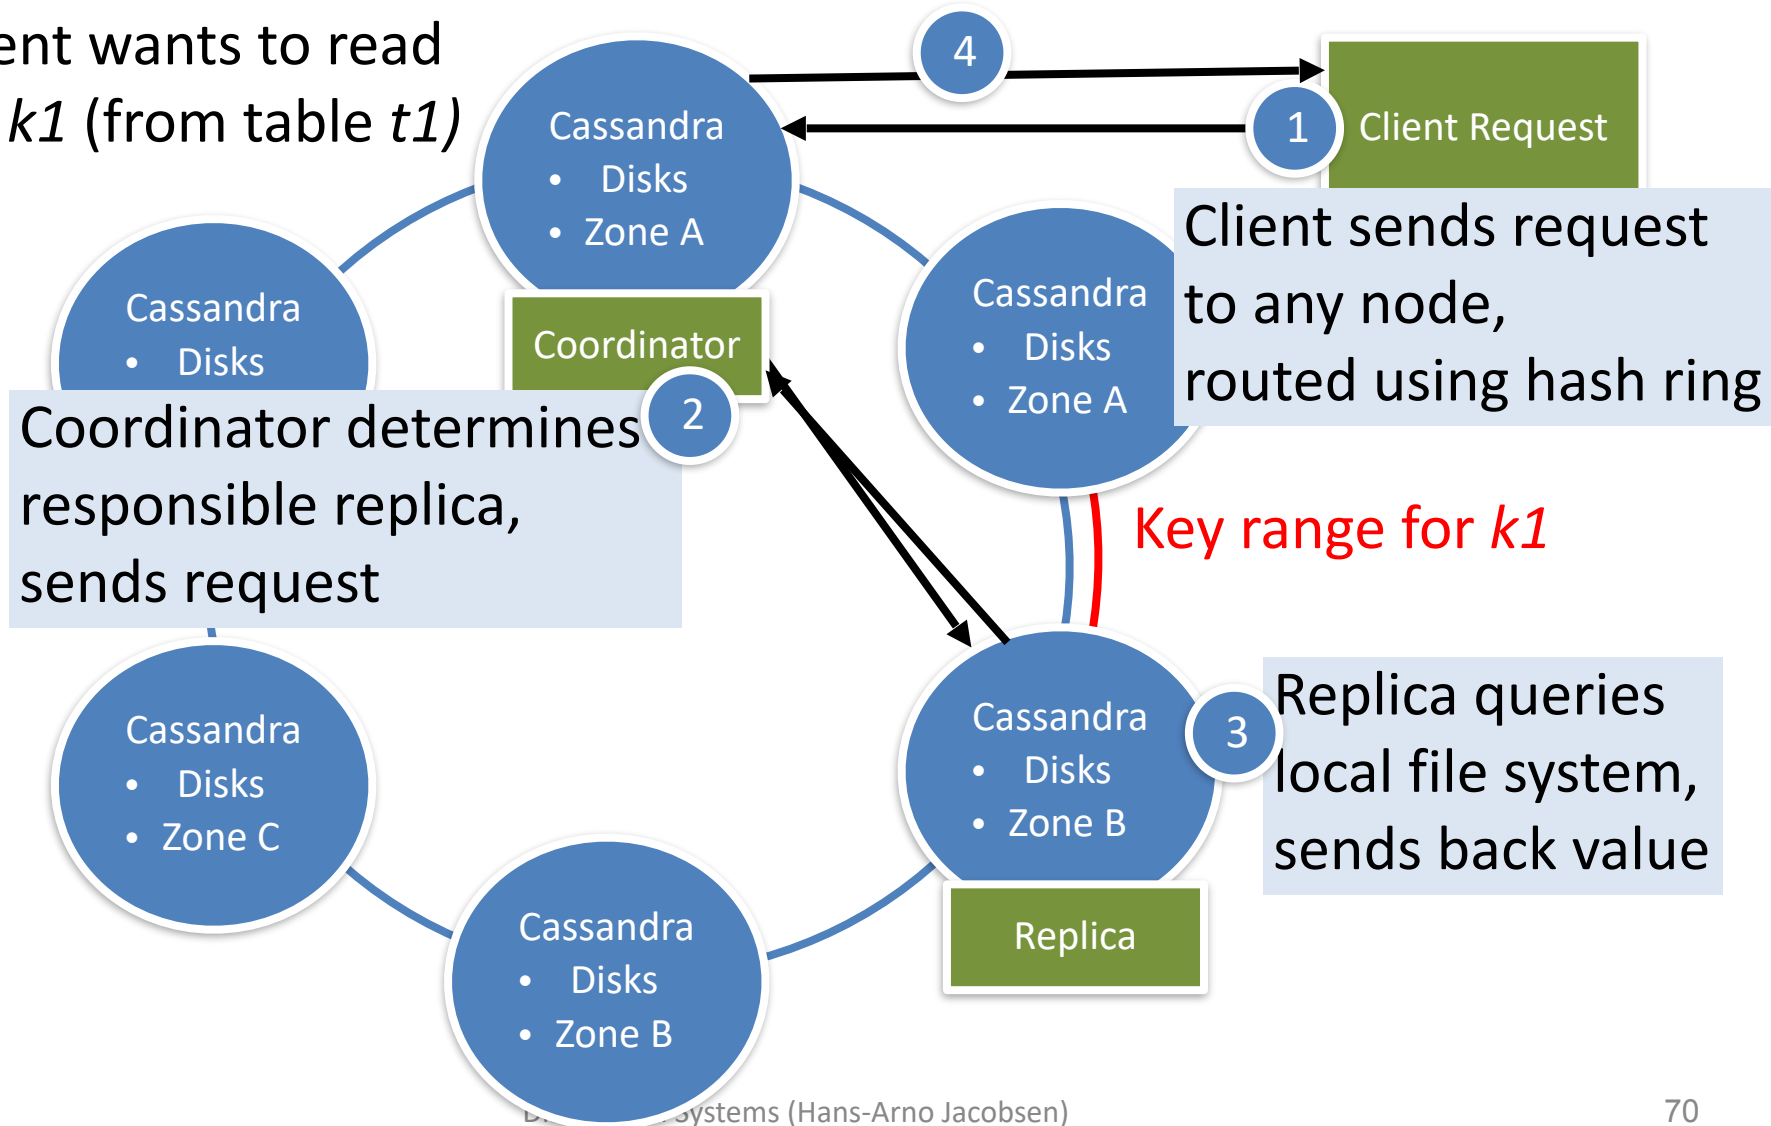
Cassandra global read-path

Client wants to read
key $k1$ (from table $t1$)



Cassandra global read-path

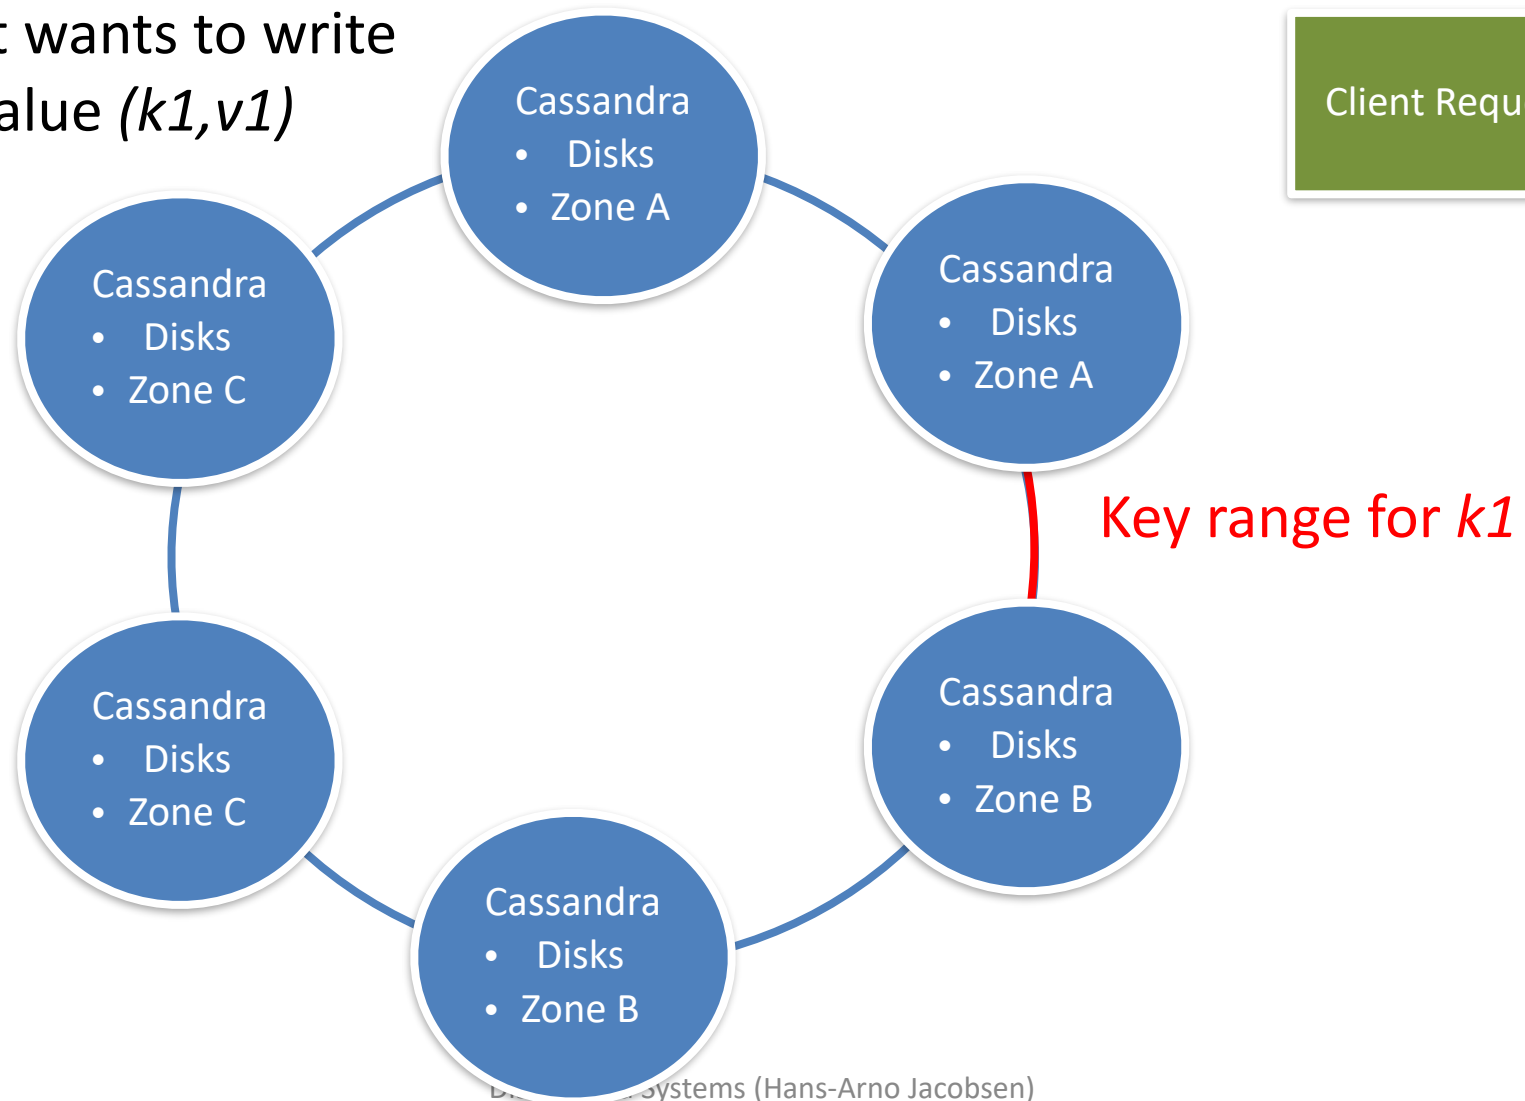
Client wants to read
key $k1$ (from table $t1$)



Cassandra global write-path

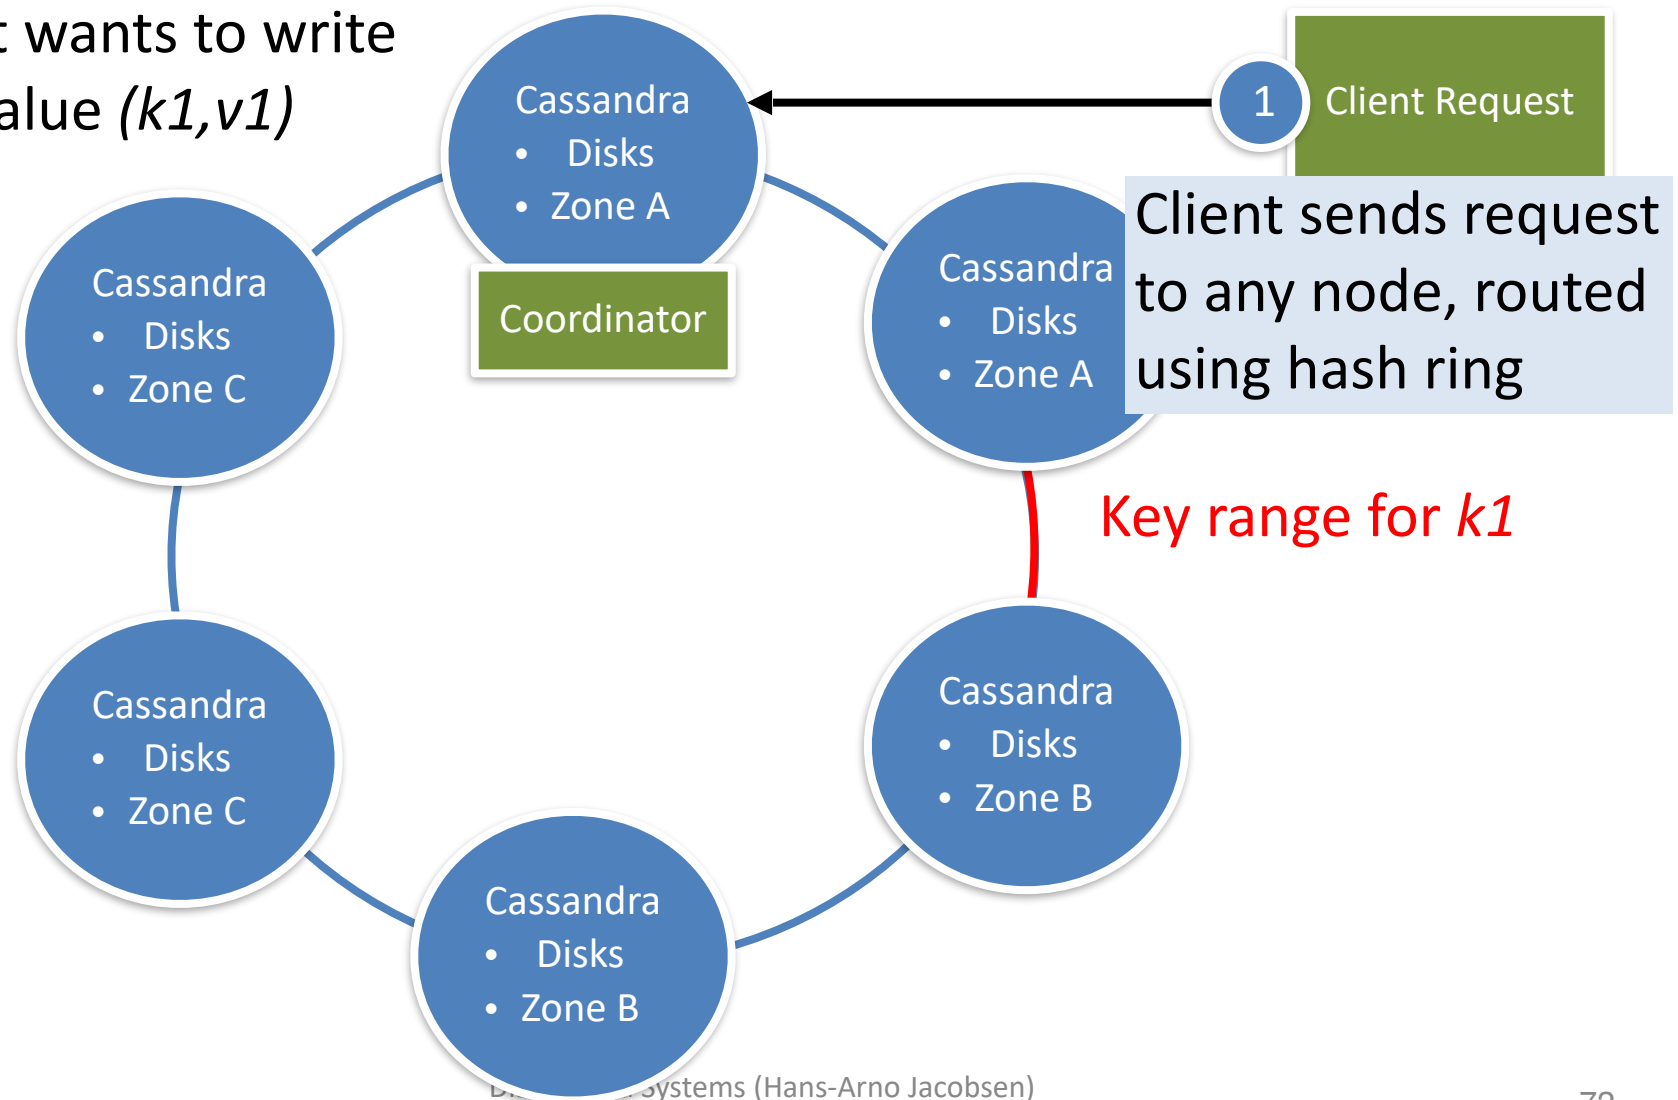
Client wants to write
key-value ($k1, v1$)

Client Request



Cassandra global write-path

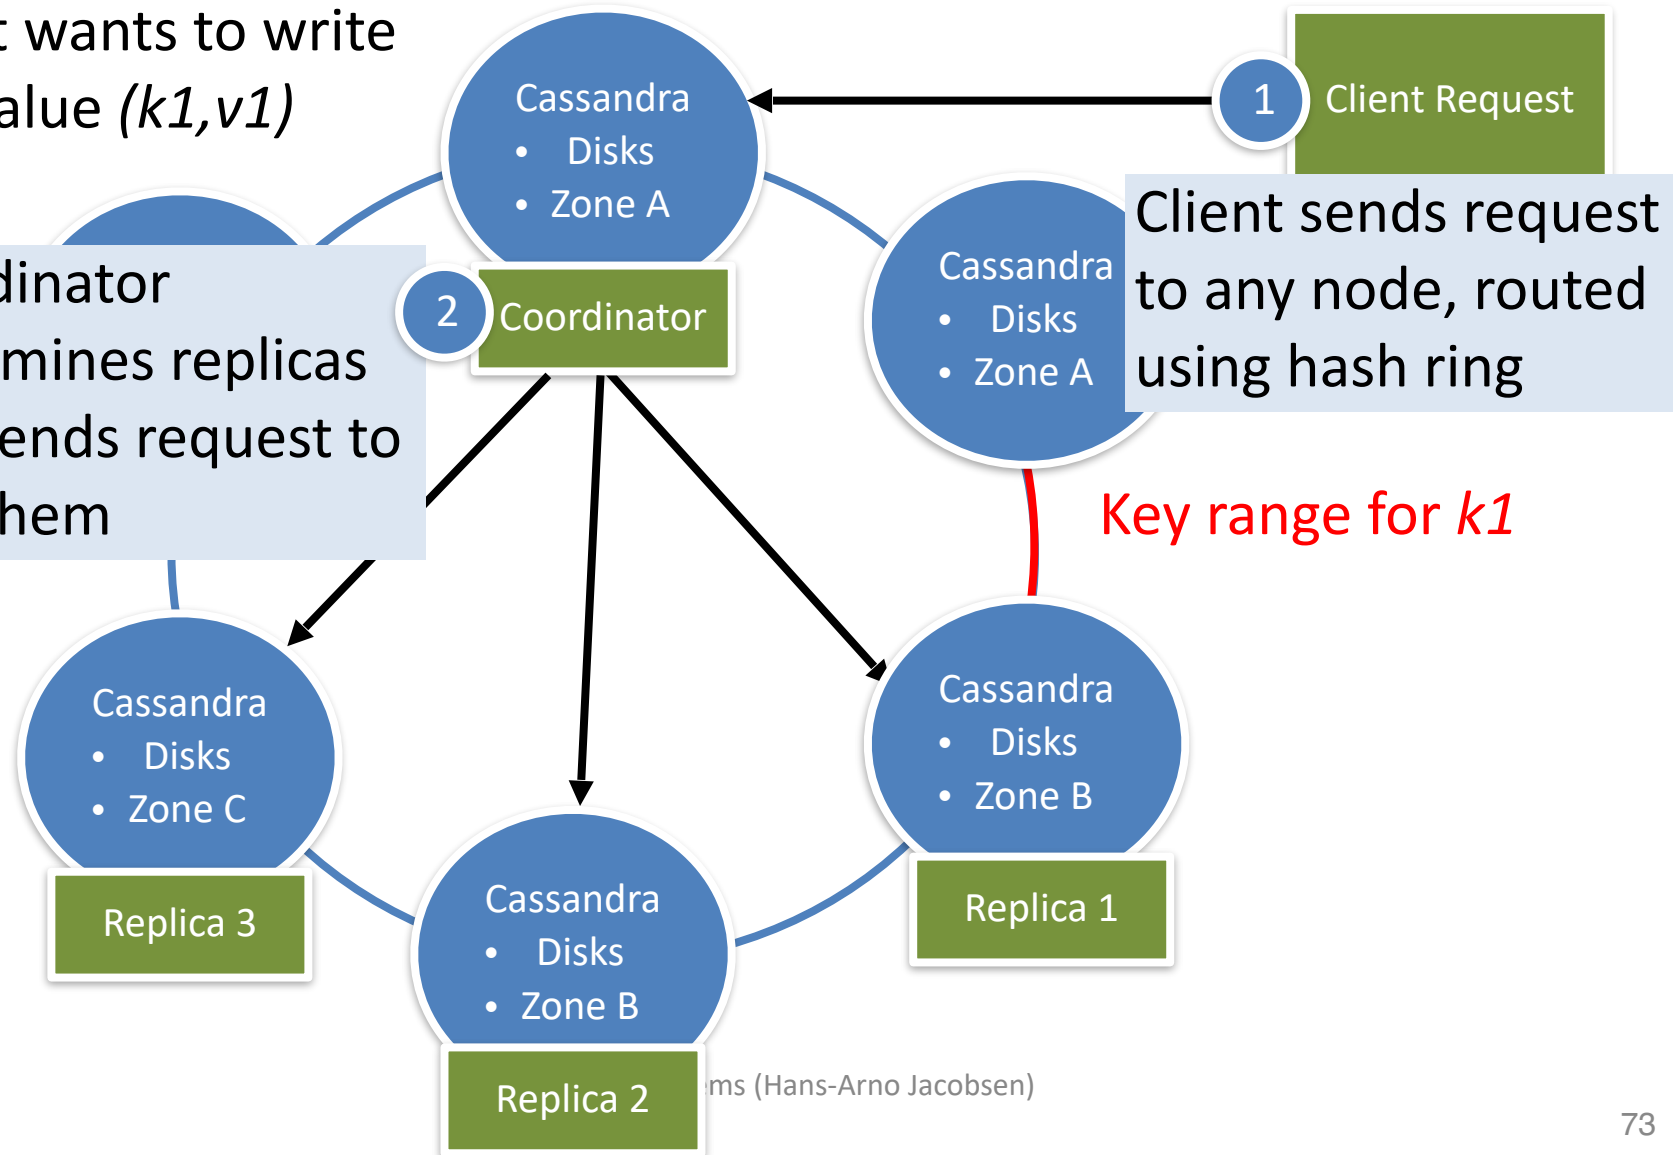
Client wants to write
key-value ($k1, v1$)



Cassandra global write-path

Client wants to write
key-value ($k1, v1$)

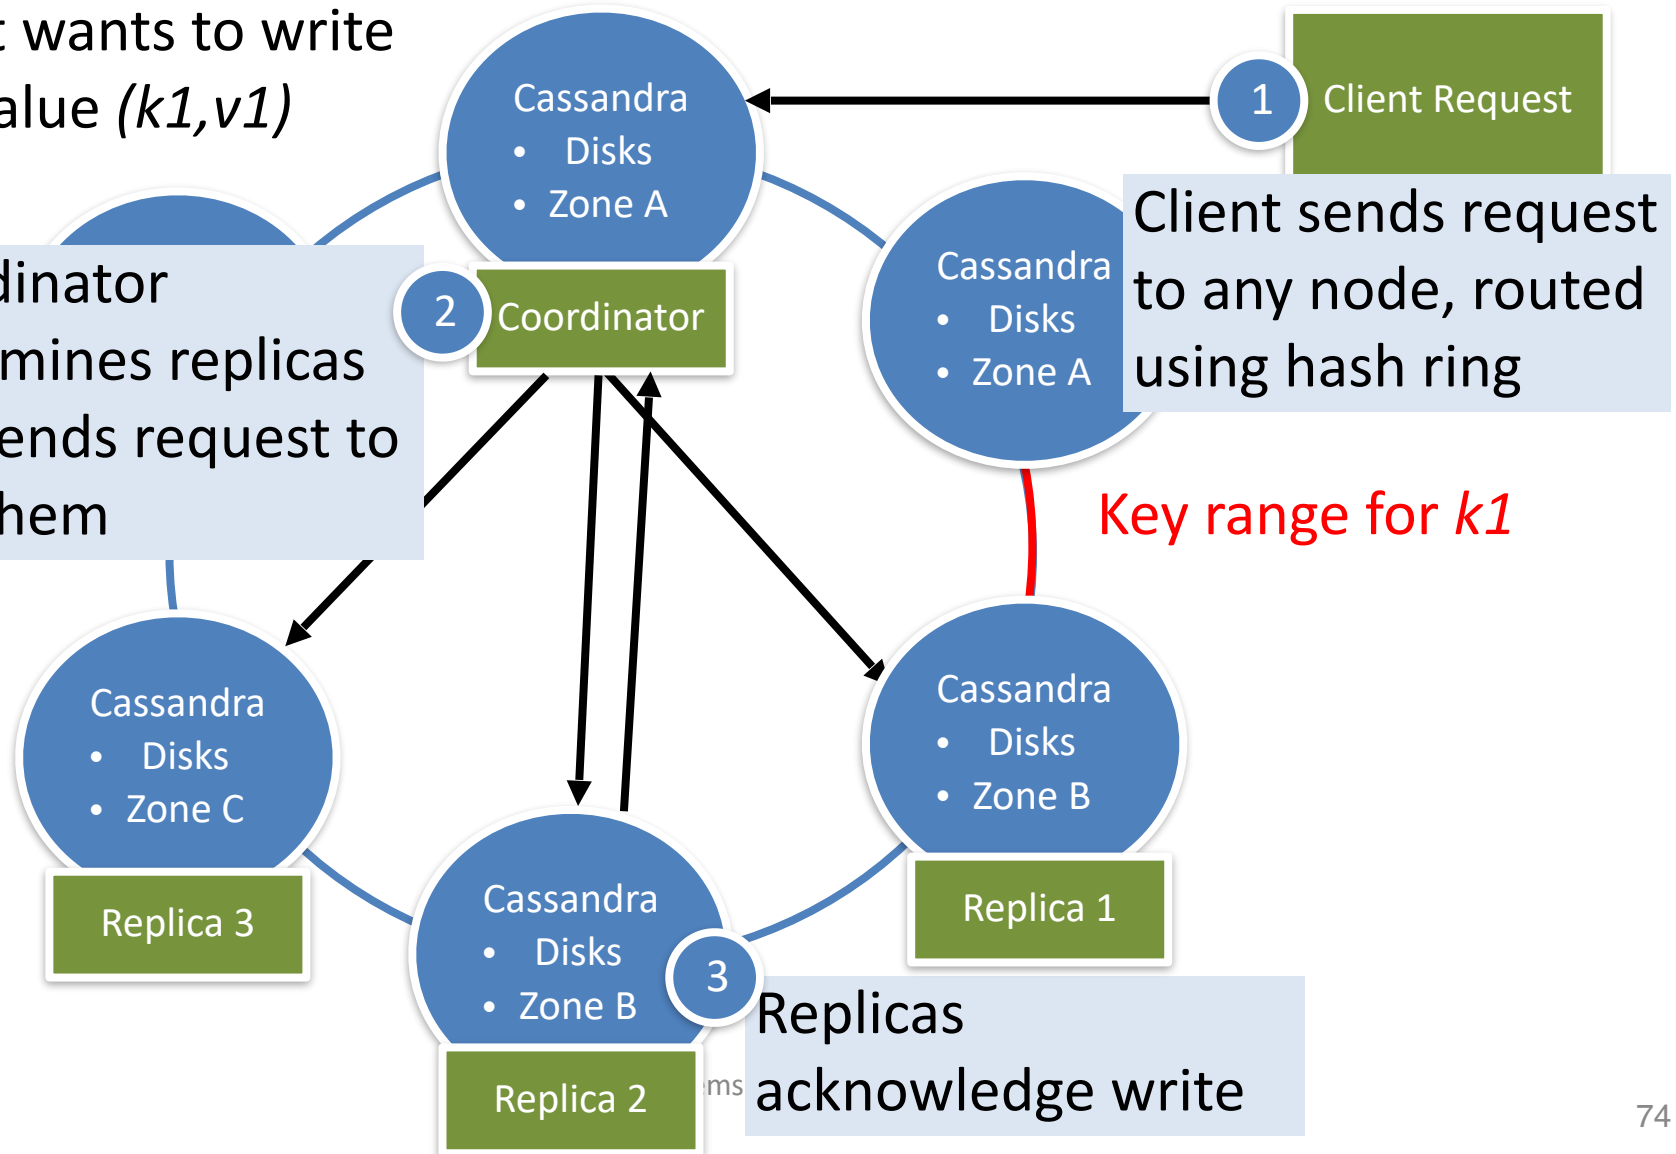
Coordinator
determines replicas
and sends request to
 n of them



Cassandra global write-path

Client wants to write
key-value ($k1, v1$)

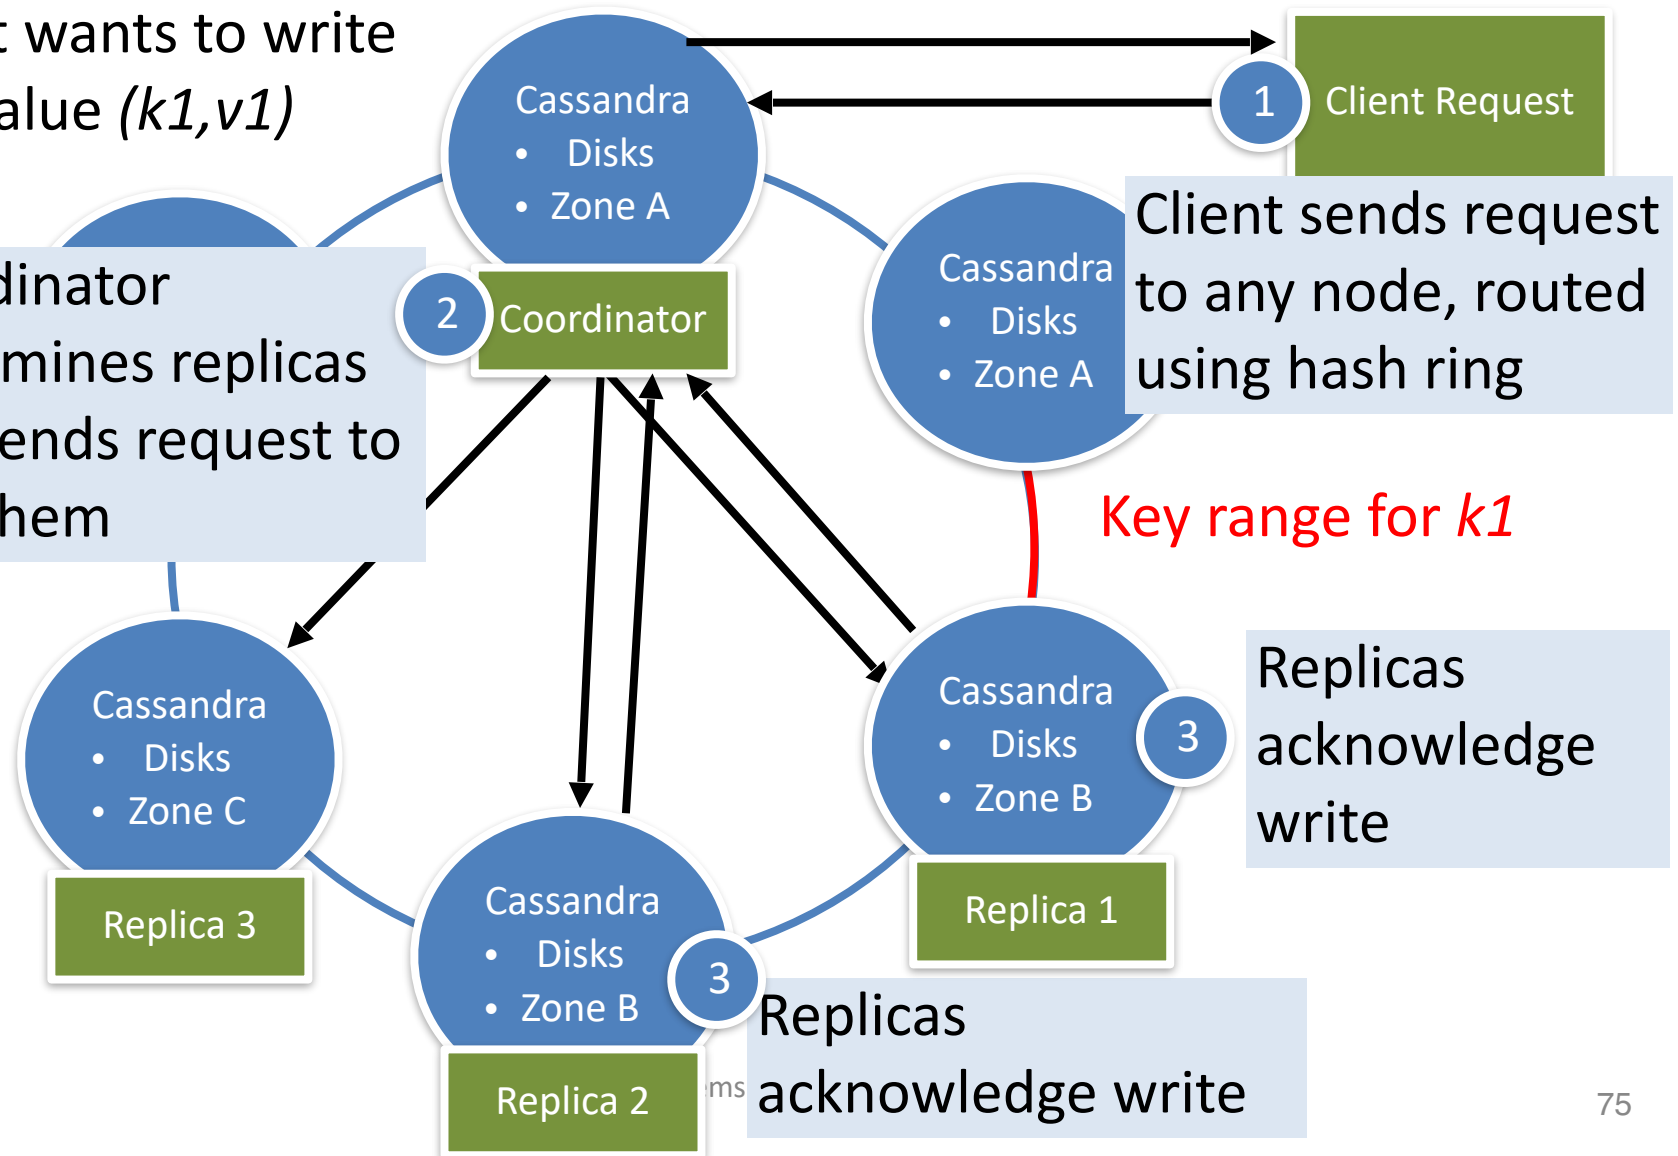
Coordinator
determines replicas
and sends request to
 n of them



Cassandra global write-path


Client wants to write key-value ($k1, v1$)

Coordinator determines replicas and sends request to n of them



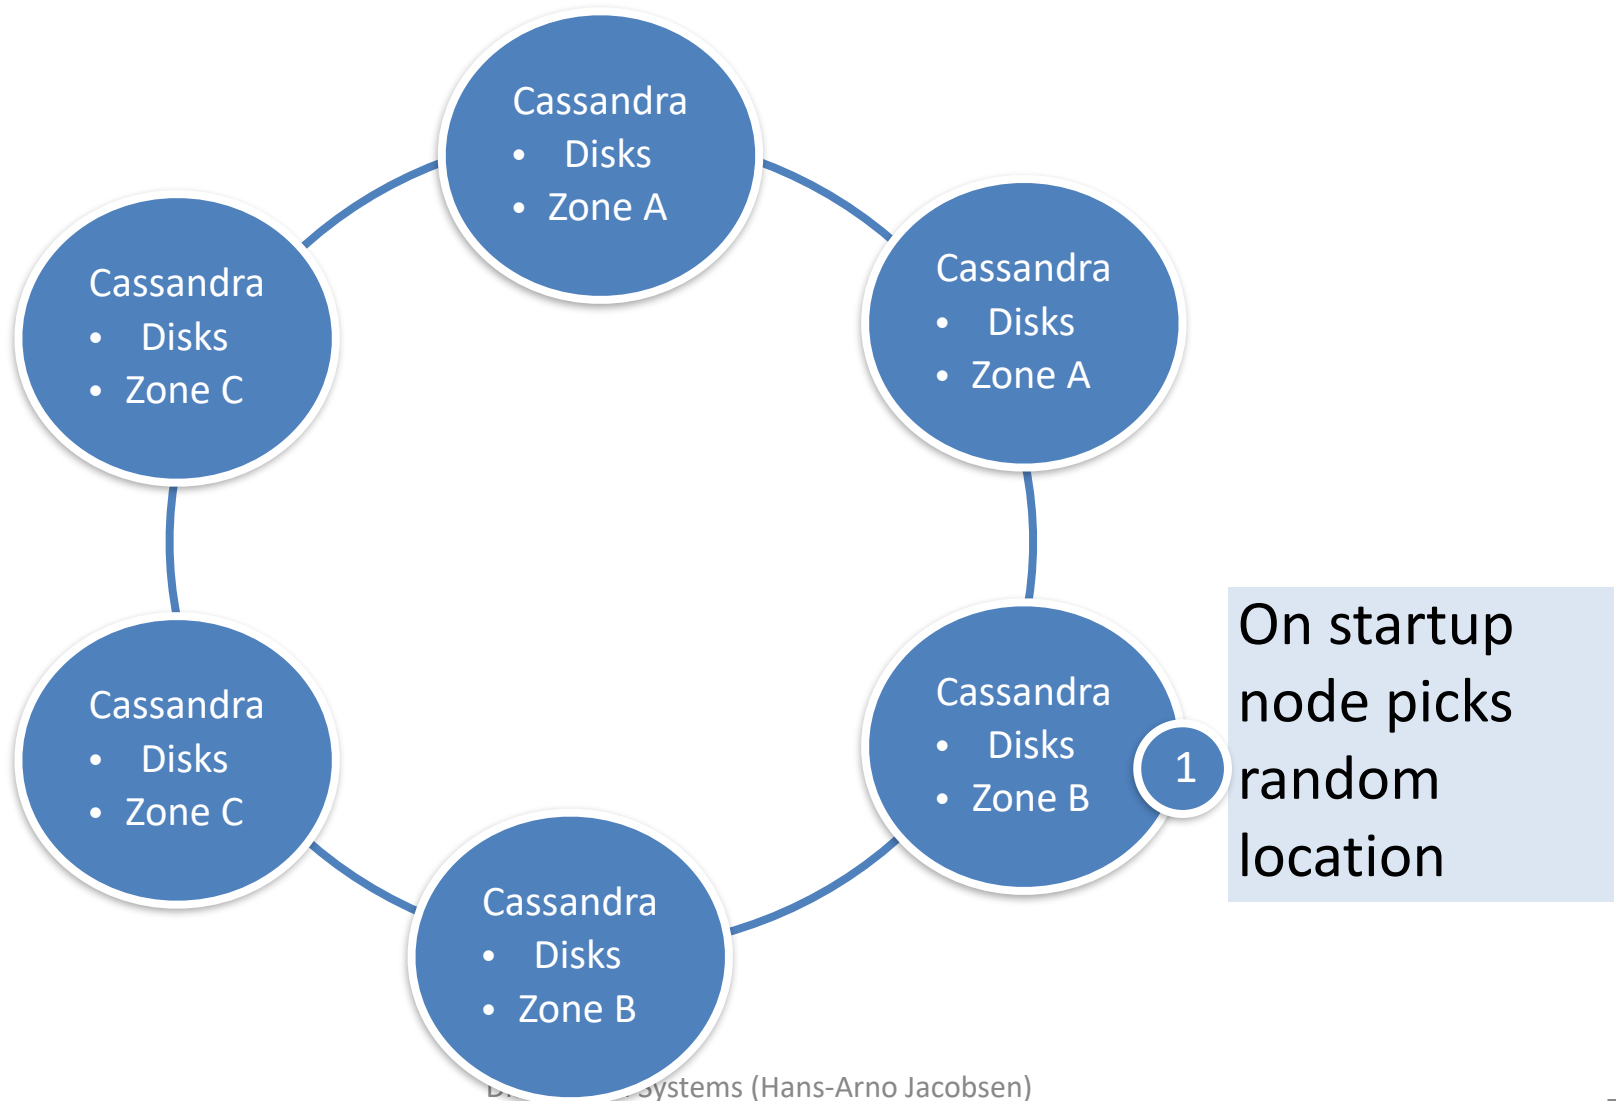
Incremental scaling in Cassandra

(i.e., adding a storage unit)



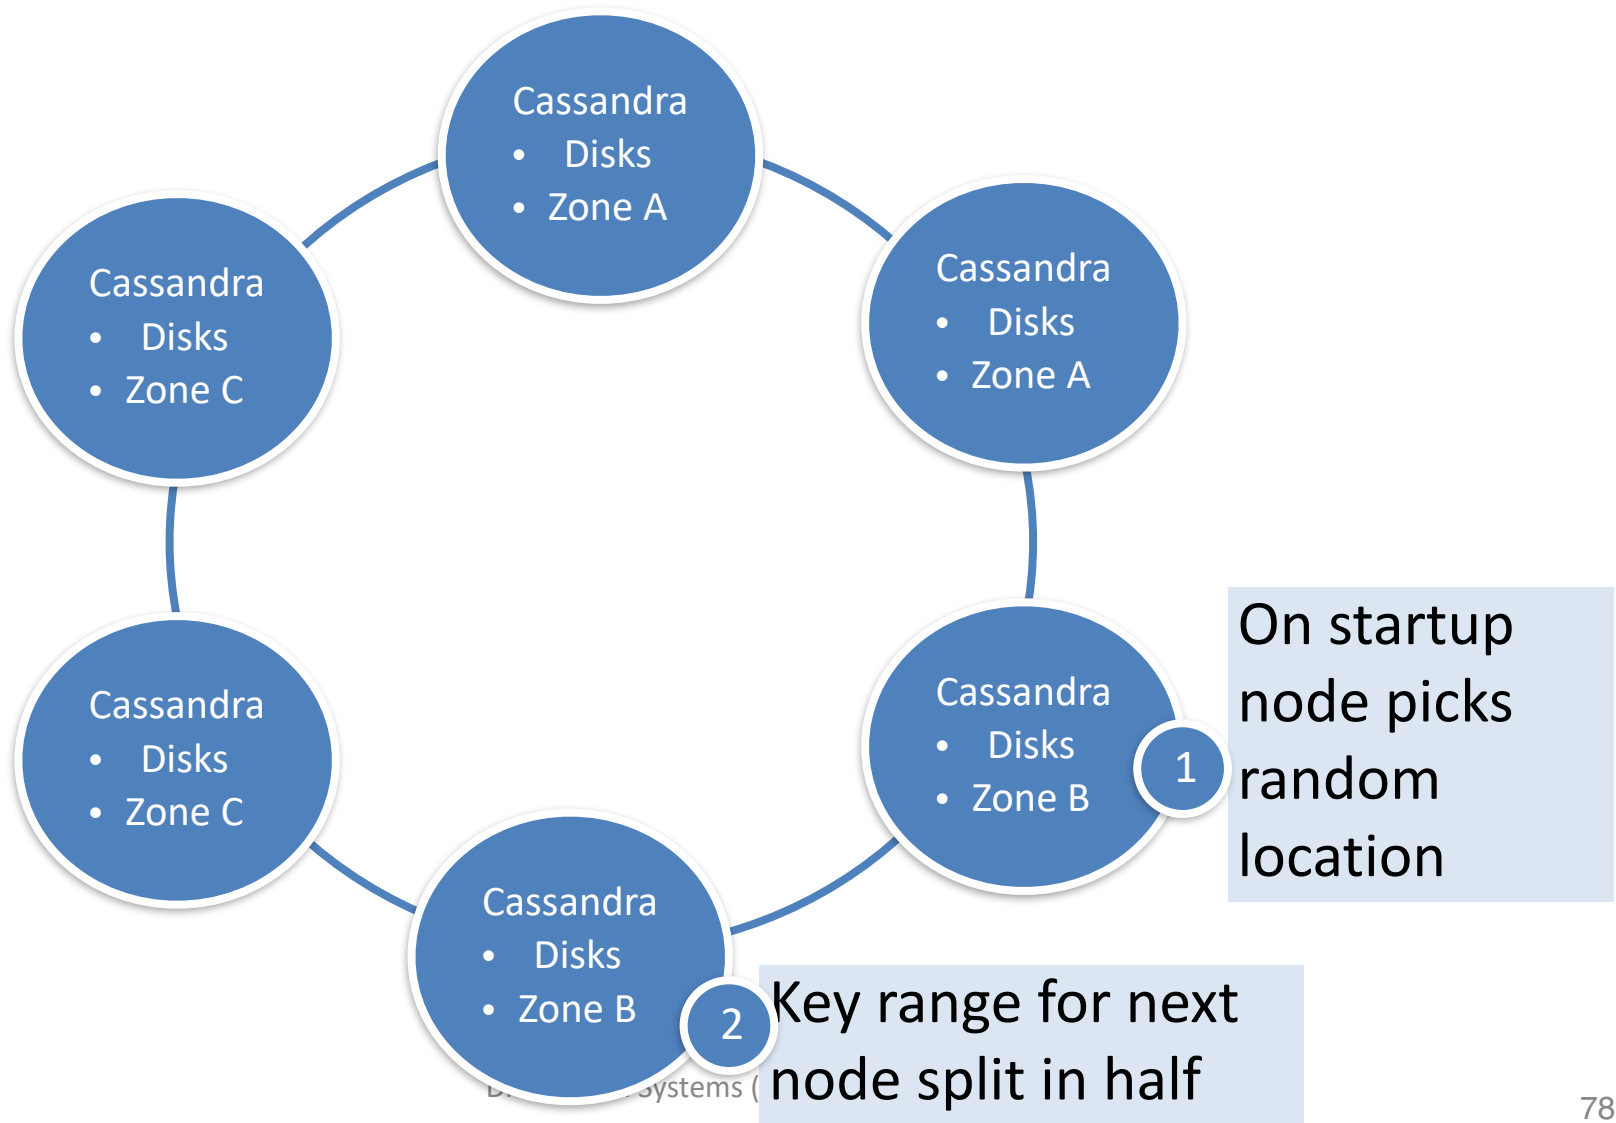
Incremental scaling in Cassandra

(i.e., adding a storage unit)



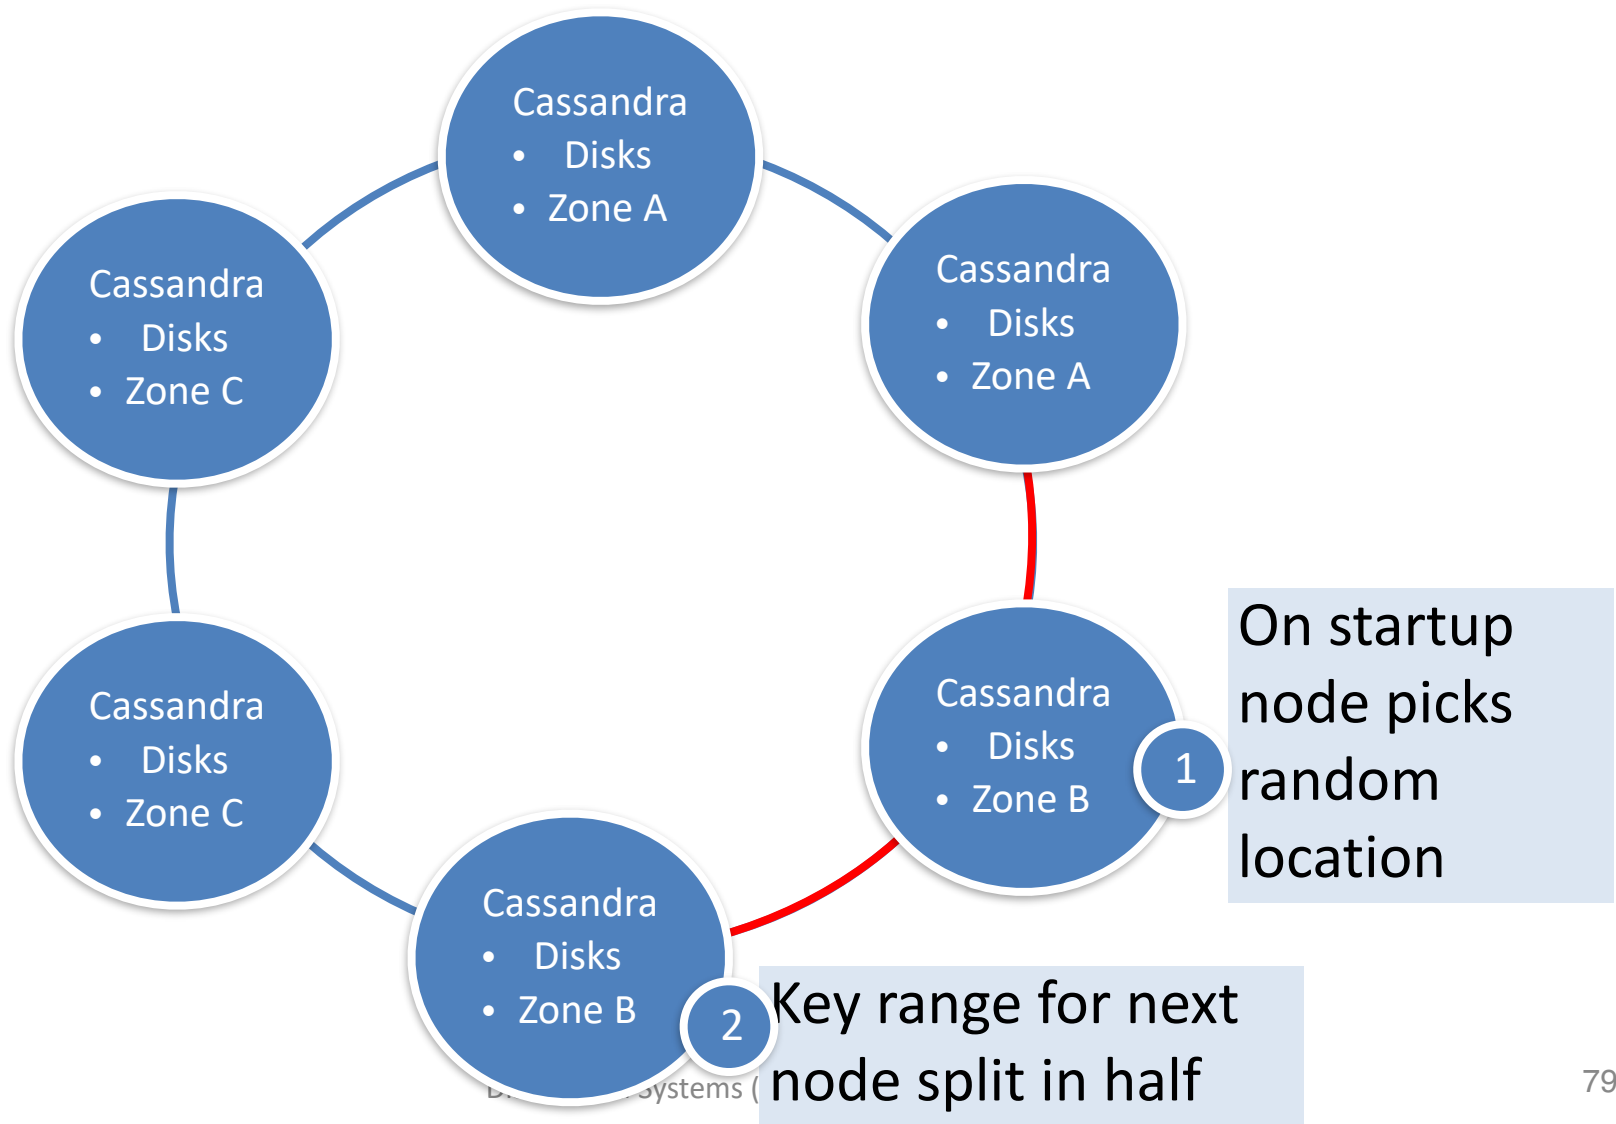
Incremental scaling in Cassandra

(i.e., adding a storage unit)



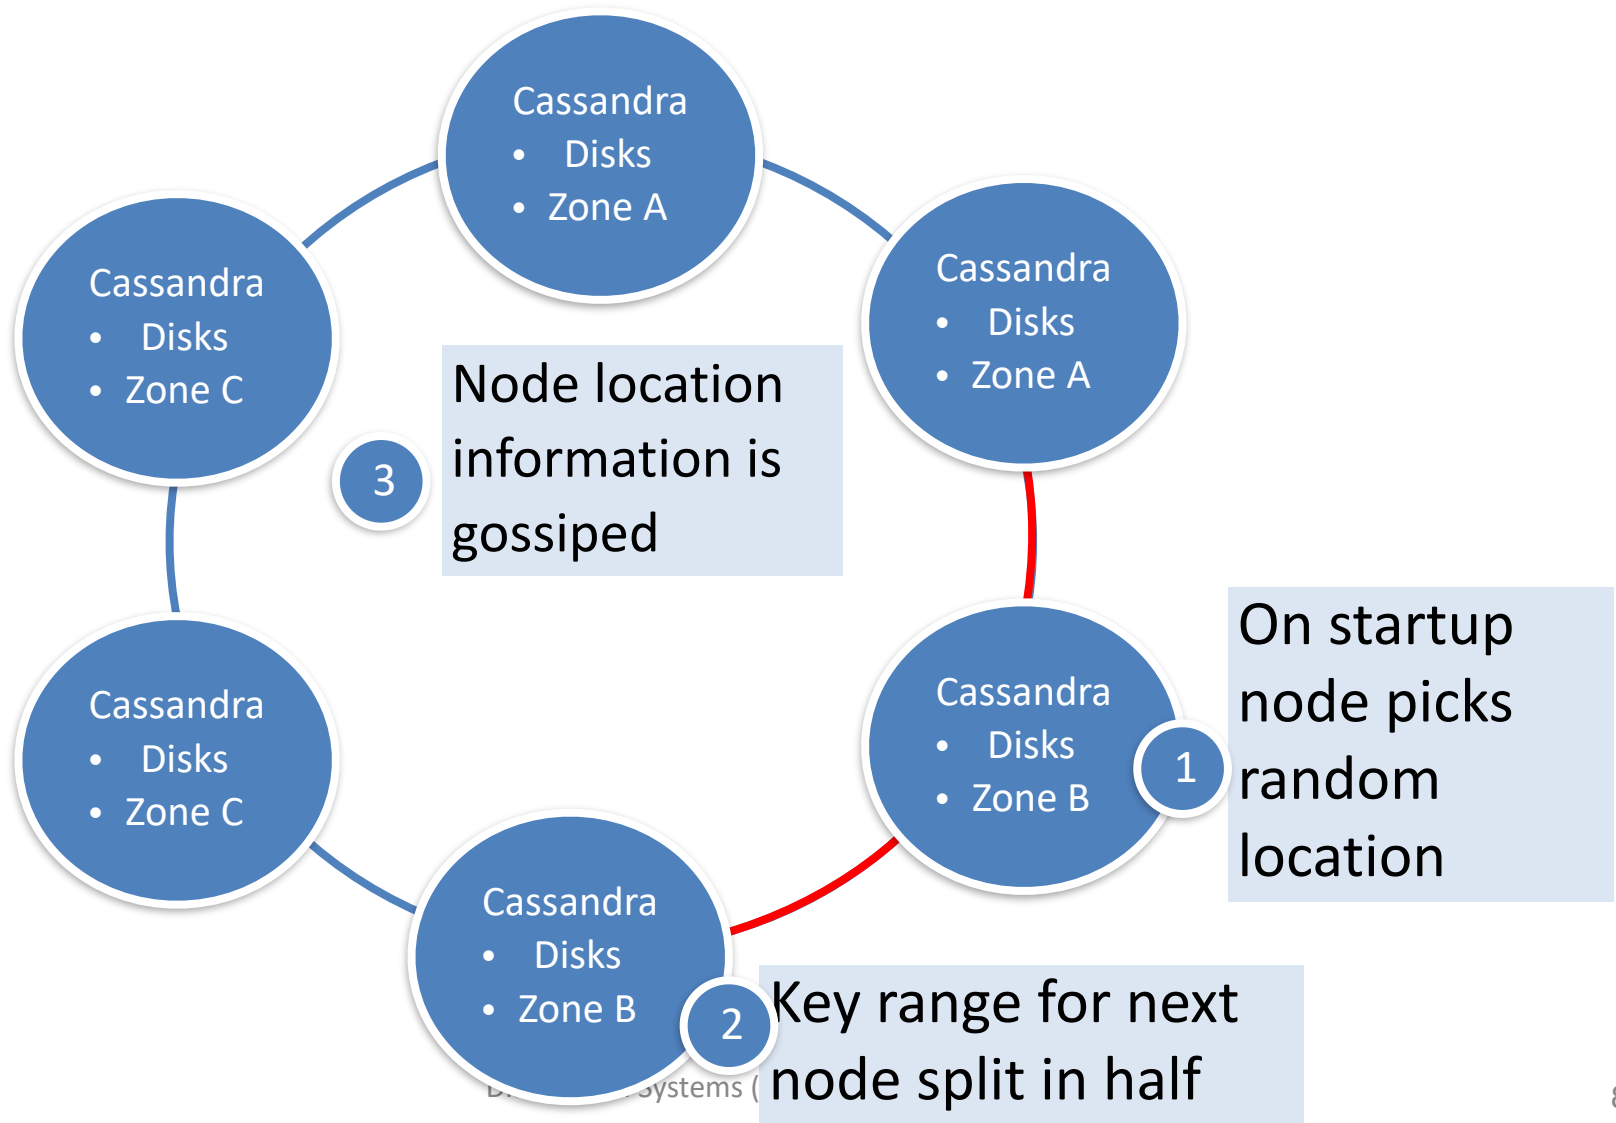
Incremental scaling in Cassandra

(i.e., adding a storage unit)



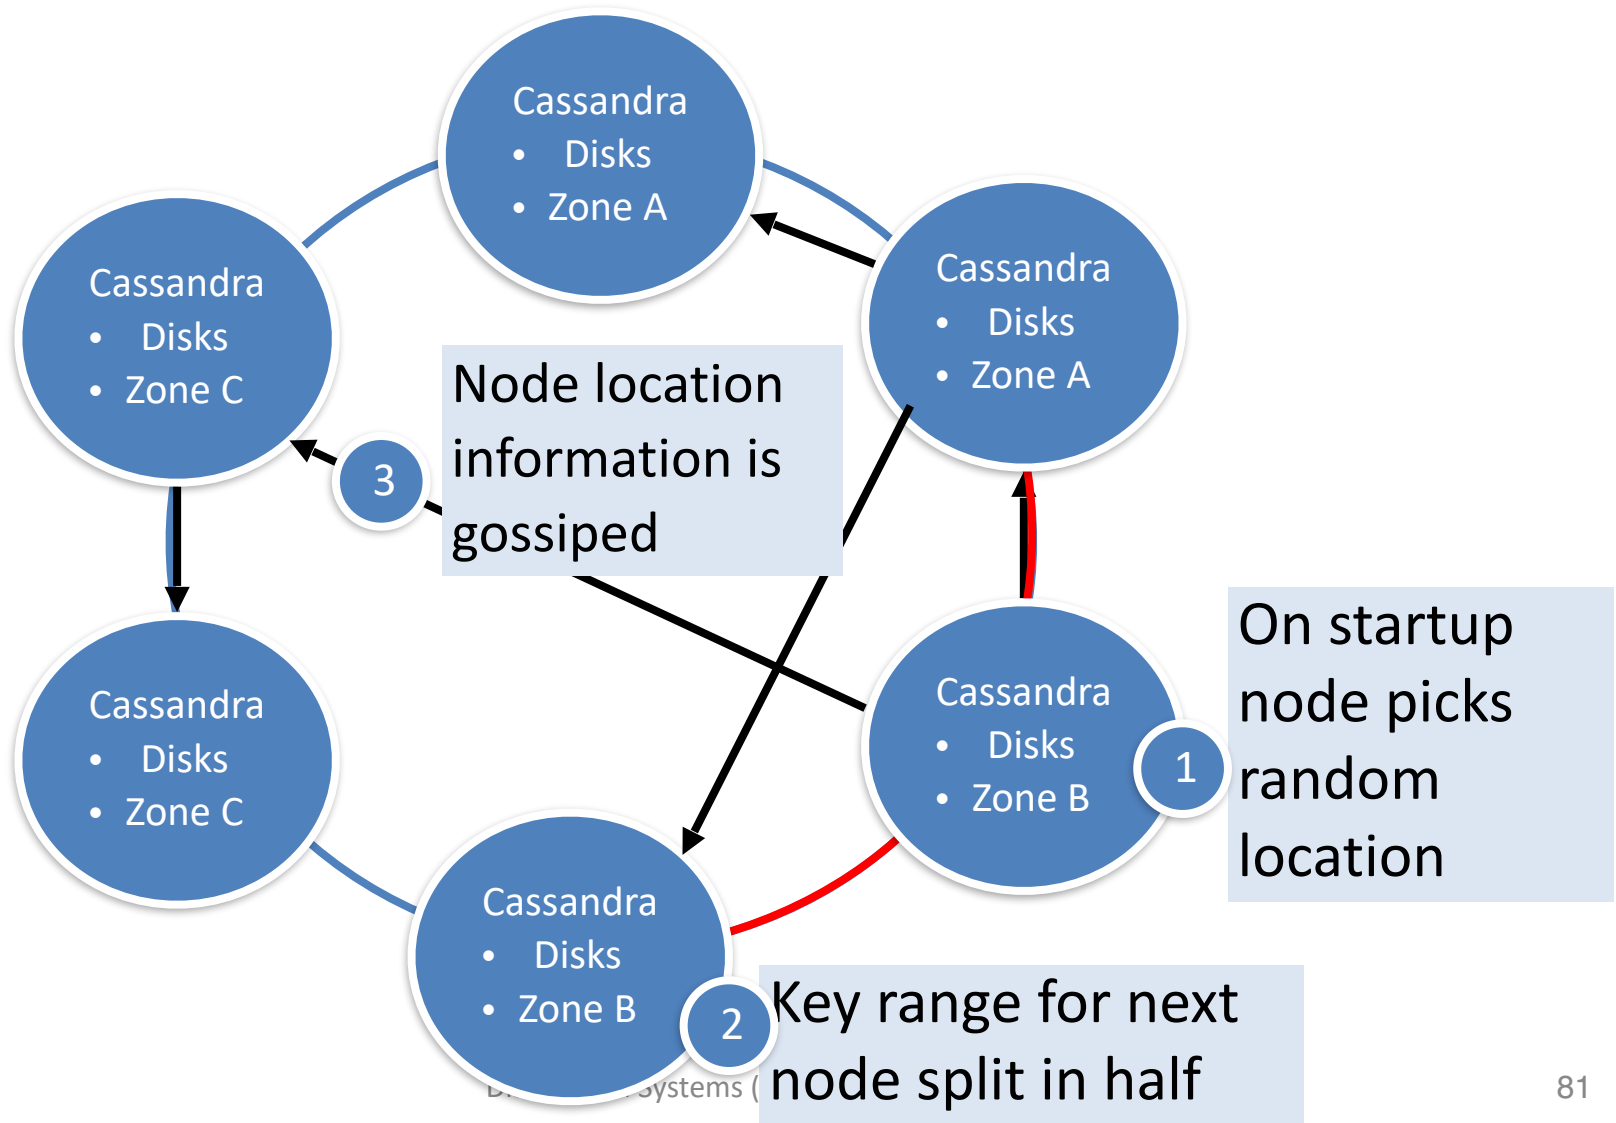
Incremental scaling in Cassandra

(i.e., adding a storage unit)

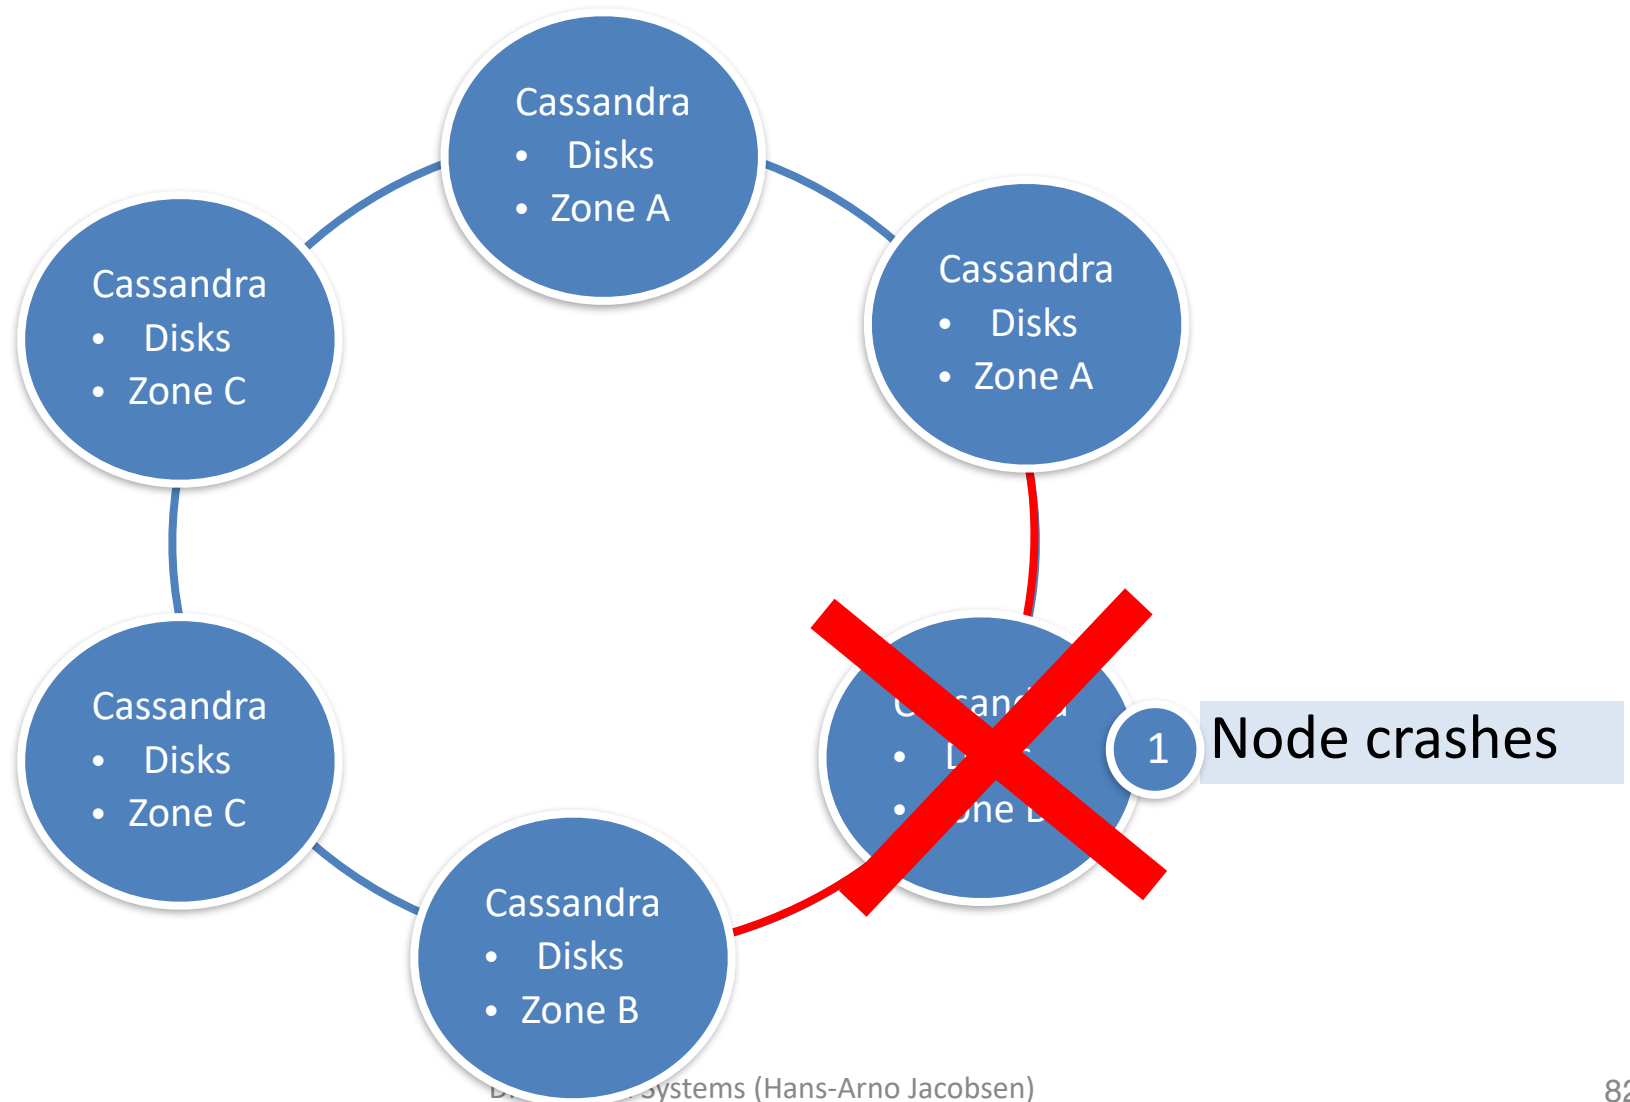


Incremental scaling in Cassandra


(i.e., adding a storage unit)



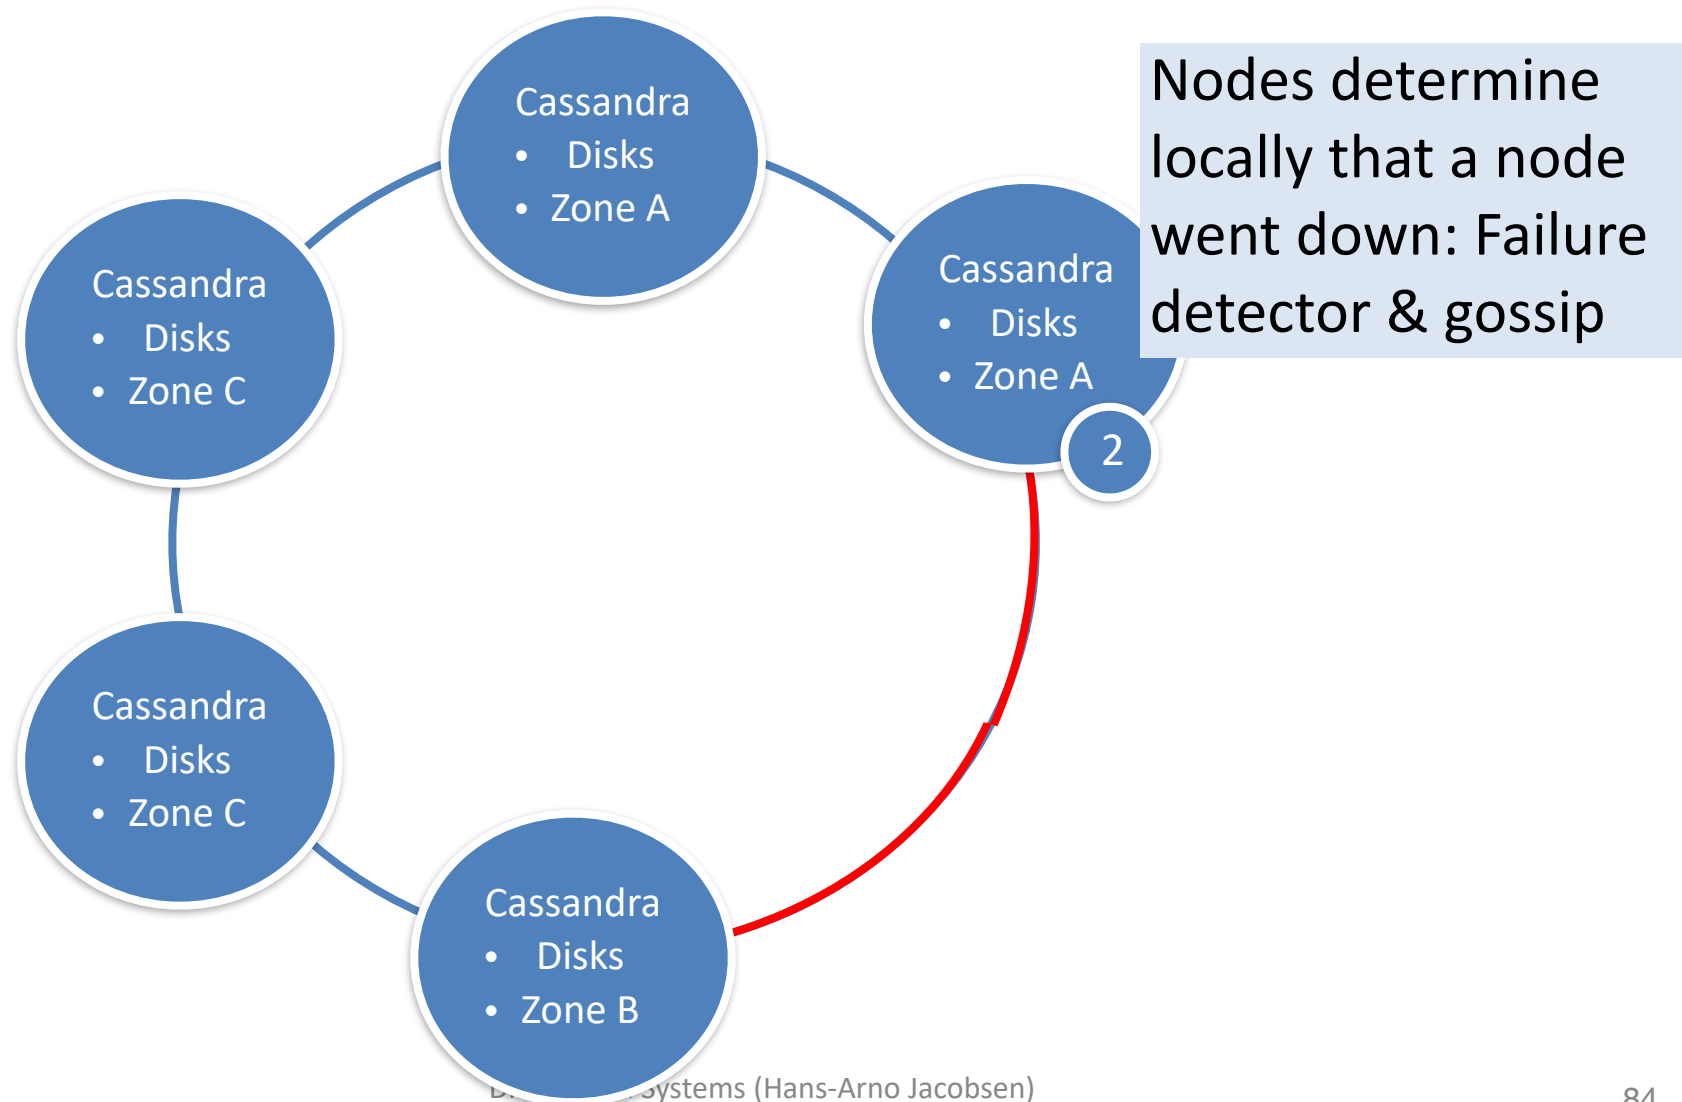
Storage unit failure



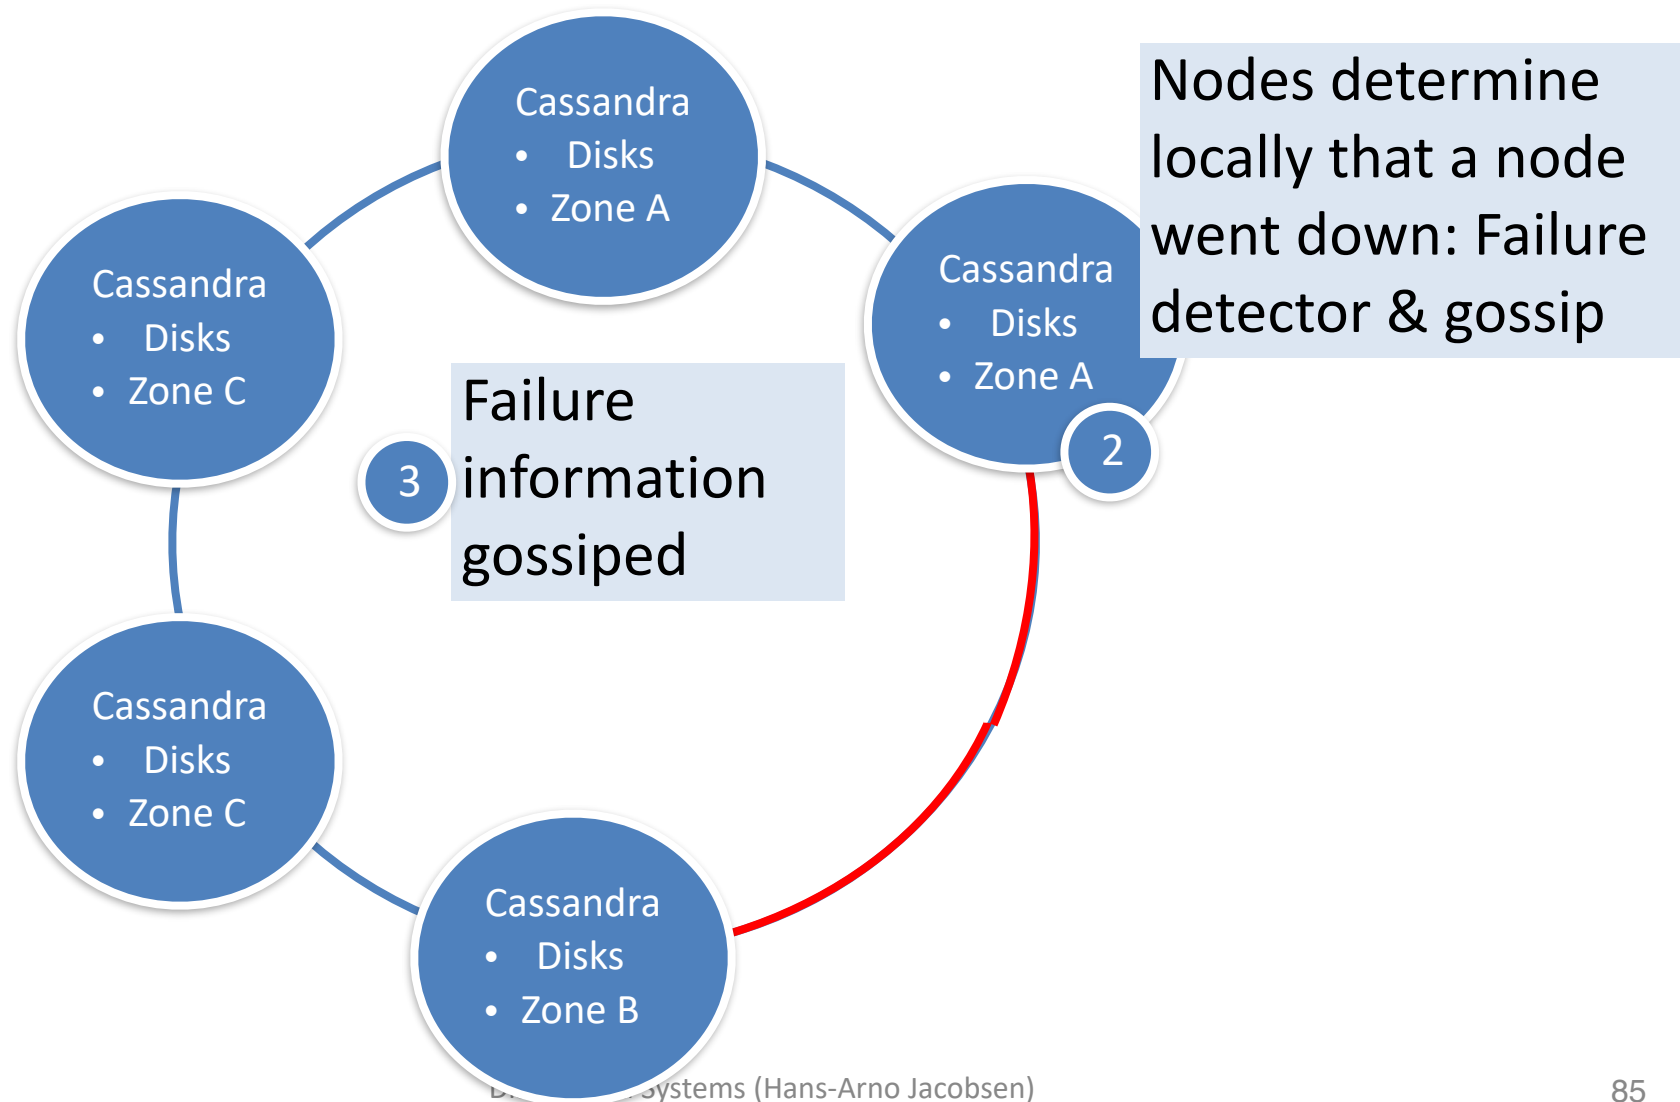
Storage unit failure



Storage unit failure



Storage unit failure



Core mechanisms

- Decentralized load balancing and scalability
 - *Cf. Consistent Hashing*
- Read/write reliability
 - *Cf. Replication*
- Membership management
 - *Cf. Gossip in Replication*
- Eventual consistency model
 - *Cf. Consistency*

Recommended Reading Materials

1. D. Karger, et al. Consistent hashing and random trees: Distributed caching protocols for relieving hot spots on the World Wide Web. In Proceedings of the Twenty-Ninth Annual ACM Symposium on Theory of Computing, pages 654-663 , 1997.

2. Cassandra by example (slides):

<https://de.slideshare.net/grro/cassandra-by-example-the-path-of-read-and-write-requests>