

# Leases in Distributed Systems

**Prof. Pramod Bhatotia**

Chair of Decentralized Systems Engineering (DSE)

Department of Computer Science

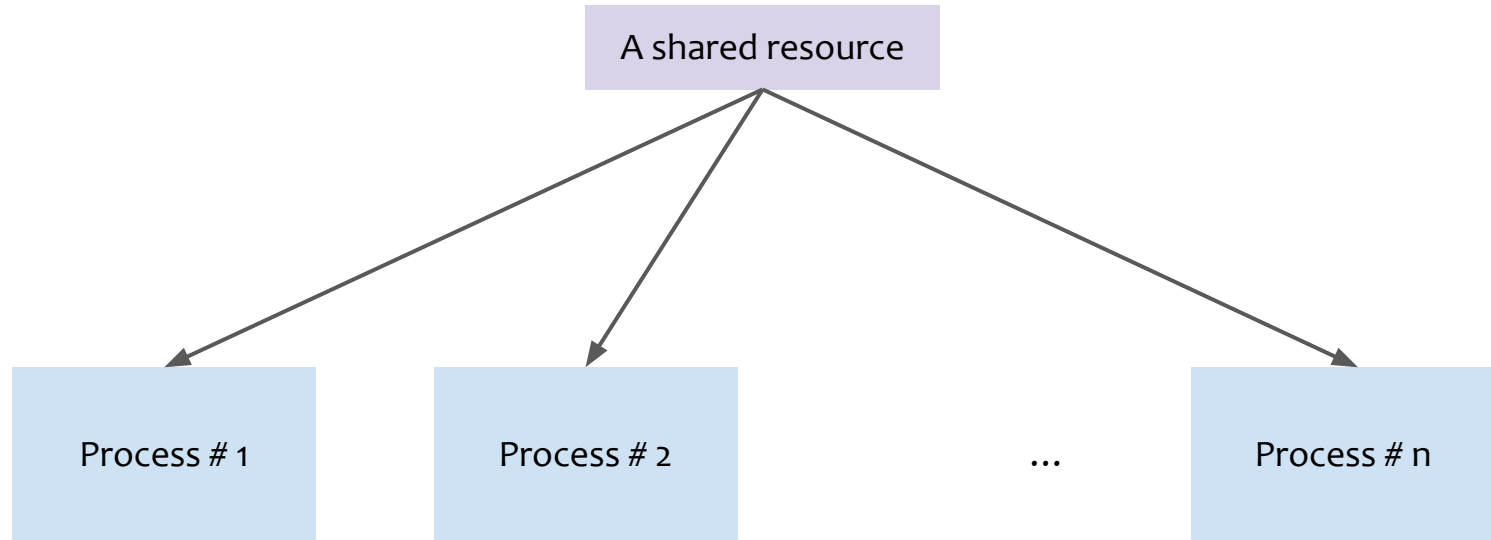
<https://dse.in.tum.de/>



# Lease

A fundamental abstraction for scalable resource management

# A typical problem



How to access a shared resource in an exclusive manner?

# Locking as a panacea?

```
acquireLock(shared_resource);
```

```
//access the shared resource
```

```
releaseLock(shared_resource);
```

- A process can crash

OR

- They can have network failure

**The problem:** A process can hold the resource indefinitely!

# Lease: An alternative fault-tolerant abstraction



- A lease is a contract that gives its holder specified rights over a resource for a limited period of time
  - A process can ask for a lease for a limited period of time, after which it expires
  - The process can renew the lease before it expires if it wants to extend the access

# Lease variants

- **Read lease** allows client to cache clean data
  - **Guarantee:** no other client is modifying data
- **Write lease** allows safe delayed writes
  - Client can locally modify the data, and then writes in a batch
  - **Guarantee:** no other client has data caches

# Using leases

- Client requests a lease
  - Specifies the resource id
- Server determines if lease can be granted
  - Read leases may be granted concurrently
  - Write leases are granted exclusively
- If conflict exists, server may send eviction notices
  - Evicted write lease must write back
  - Evicted read leases must flush/disable caching
  - Client acknowledges when completed

# Bounded lease term simplifies recovery

- Before lease expires, client must renew lease
- Client fails while holding a lease?
  - Server waits until the lease expires, then unilaterally reclaims
  - If client fails during eviction, server waits then reclaims
- Server fails while leases outstanding? On recovery:
  - Wait lease period + clock skew before issuing new leases



- Leases depend on well-behaved clocks
  - If the clock speeds are mismatched between the lease grantor and the lease holder, it can lead to inconsistencies

# More example use-cases

- Google's Chubby and Apache Kafka: Co-ordination services
  - <https://research.google/pubs/pub27897/>
  - <https://zookeeper.apache.org/>
- Apache Kafka: an event streaming platform
  - <https://kafka.apache.org/>
- Google's Thialfi – a client notification service
  - <https://research.google/pubs/pub37474/>
- DHCP protocol
  - [https://en.wikipedia.org/wiki/Dynamic\\_Host\\_Configuration\\_Protocol](https://en.wikipedia.org/wiki/Dynamic_Host_Configuration_Protocol)
- Google's Slicer: Auto-sharding
  - <https://research.google/pubs/pub46921/>

- **Seminal paper:**

- <https://web.stanford.edu/class/cs240/readings/89-leases.pdf>

## Leases: An Efficient Fault-Tolerant Mechanism for Distributed File Cache Consistency

Cary G. Gray and David R. Cheriton  
Computer Science Department  
Stanford University