

Background on Microkernels

Prof. Pamod Bhatotia

<https://dse.in.tum.de/bhatotia/>



Monolithic Kernels

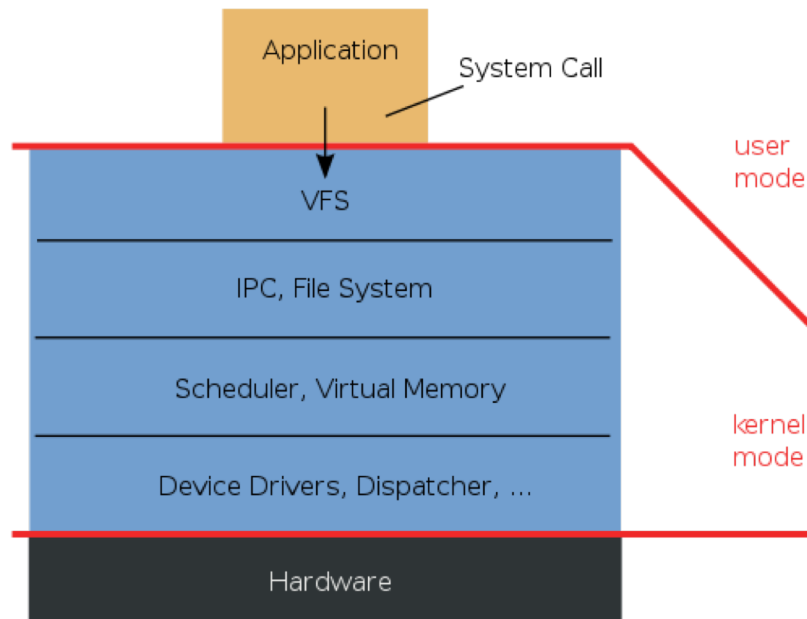
- All OS services operate in kernel space
- Good performance
- Disadvantages
 - Dependencies between system component
 - Complex & huge (millions(!) of lines of code)
 - Larger size makes it hard to maintain
- E.g. Multics, Unix, BSD, Linux

Microkernels

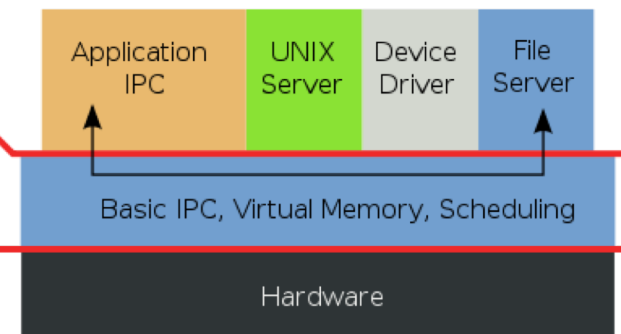
- Minimalist approach
 - IPC, virtual memory, thread scheduling
- Put the rest into user space
 - Device drivers, networking, file system, user interface
- More stable with less services in kernel space
- Disadvantages
 - Lots of system calls and context switches
- E.g. Mach, L4, AmigaOS, Minix, K42

Microkernel Vs Monolithic Kernels

Monolithic Kernel based Operating System



Microkernel based Operating System



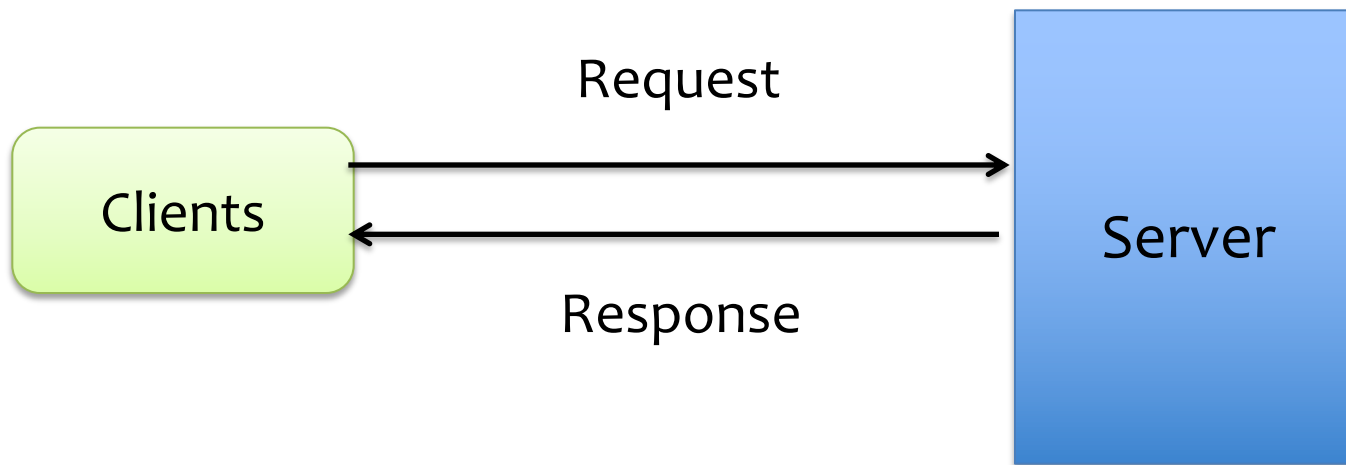
Background on State Machine Replication

Prof. Pamod Bhatotia


<https://dse.in.tum.de/bhatotia/>



Fault tolerance



Replication



State machine replication

1. Implement a service as a deterministic state machine
2. Replicate server
3. Provide all replicas with the same input

Guarantees: all correct replicas will produce the same output

Consensus algorithm

Distributed Synchronization via Zookeeper

Prof. Pamod Bhatotia

<https://dse.in.tum.de/bhatotia/>



Co-ordination in distributed systems



- (Dynamic) configuration
- Synchronization
- Leader election
- Group membership
- Barriers
- Locks
- ...

Challenges in distributed systems



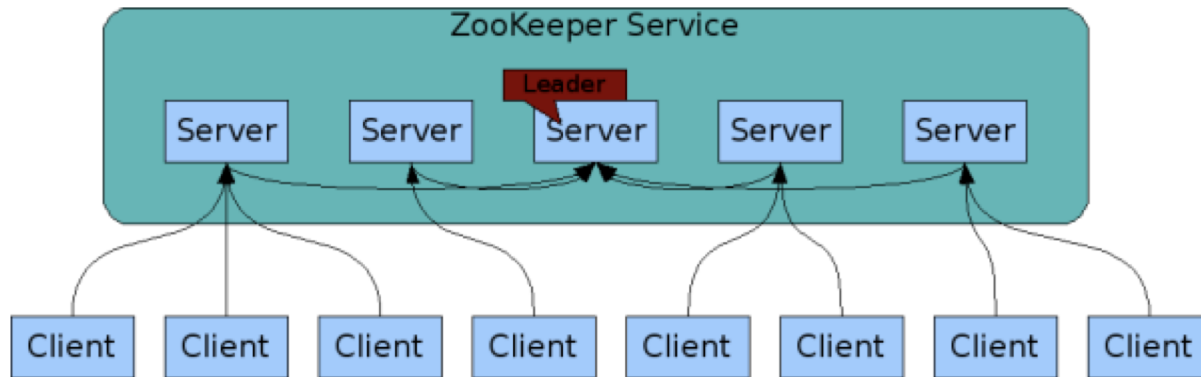
- The network is unreliable
- Process may crash/fail in arbitrary ways
- The network messages may arrive arbitrarily

- A co-ordination (micro-)kernel
 - Minimalistic APIs that can be used to build a wide-range of co-ordination primitives
- APIs are wait-free
 - No blocking primitives in ZooKeeper
 - Blocking can be implemented by a client
 - Deadlock free!

Zookeeper design principles

- Zookeeper = FIFO ordering for clients + Linearizable writes + Wait-free APIs
- Guarantees
 - Client requests are processed in FIFO order
 - Writes to ZooKeeper are linearizable
 - Clients receive notifications of changes before the changed data becomes visible

Zookeeper architecture




Terminology:

- **Clients:** users of the Zookeeper service
- **Server:** process providing the Zookeeper service
- **Session:** Clients establish a session when connecting to Zookeeper

Zookeeper data model

Abstraction: A set of data nodes (znodes) organized in a hierarchal namespace



Znodes can store data

znodes

- Znodes are accessed similar to UNIX filesystem namespace
- Znodes can be classified as:
 - **Regular:** created and deleted by clients explicitly
 - **Ephemeral:** can be deleted explicitly or by the system itself when the session terminates (the client that created it)
- **Flags:**
 - **Sequential:** monotonically increasing counters
 - **Watch flag:** allows client to receive timely notification of changes without polling


Zookeeper APIs

- `create(path, data, flags)`
- `delete(path, version)`
- `exists(path, watch)`
- `getData(path, watch)`
- `setData(path, data, version)`
- `getChildren(path, watch)`
- `Sync()`

- Version is used for conditional update
- Synchronous and asynchronous APIs are available for clients!


How to use Zookeeper to implement distributed co-ordination protocols?

Example 1: Configuration



1. Workers get configuration
 - `getData("../config/settings", true)`
2. Admin change the config
 - `setData("../config/settings", newConf-1)`
3. Workers notified of change and get the new settings
 - `getData("../config/settings", true)`


Example 2: Group membership



1. Register serverName in group
 - Create(“.../workers/workerName”, hostInfo, EPHEMERAL)
2. List group members
 - getChildren(“.../workers”, true)


Example 3: Leader Election

1. `getData(.../workers/leader", true)`
2. If successful follow the leader described in the data and exit
3. `create(.../workers/leader", hostname, EPHERMERAL)`
4. If successful lead and exit
5. Goto step 1



Example 4: Locks

1. `create(.../locks/mutex", Ephemeral)`
2. If succeed then lock acquired
3. Else, `getData(.../locks/mutex", true)`
4. Goto step 1




Example 5: Locks without herding

Lock

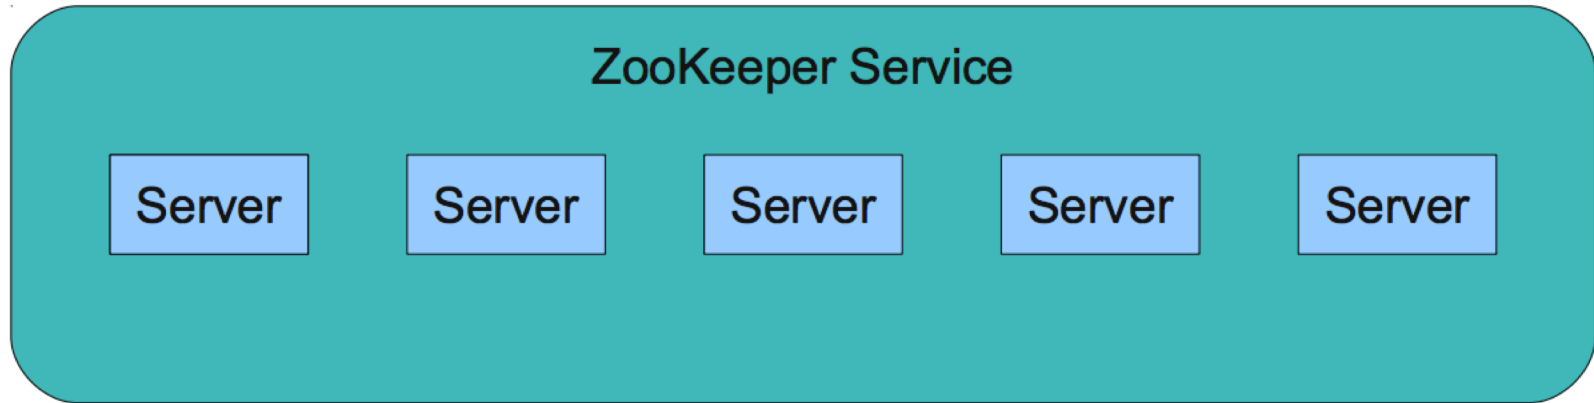
```
1 n = create(l + "/lock-", EPHMERAL|SEQUENTIAL)
2 C = getChildren(l, false)
3 if n is lowest znode in C, exit
4 p = znode in C ordered just before n
5 if exists(p, true) wait for watch event
6 goto 2
```

Unlock

```
1 delete(n)
```



System Implementation



- All servers have a copy of the state in memory
- A leader is elected at startup
- Followers service clients, all updates go through leader
- Update responses are sent when a majority of servers have persisted the change
 - We need $2f+1$ machines to tolerate f failures

Summary

- Apache Zookeeper
 - Co-ordination in distributed systems
 - A distributed co-ordination kernel
 - Usage to build powerful primitives

- Resources:
 - **Compulsory reading:** Zookeeper [ATC'10]:
 - Website: <https://zookeeper.apache.org>
 - Paper: https://www.usenix.org/legacy/event/atc10/tech/full_papers/Hunt.pdf
 - **Recommended reading:**
 - Chubby [OSDI'06]: <https://research.google.com/archive/chubby.html>
 - Zab [DSN'11]: <https://dl.acm.org/citation.cfm?id=2056409>
 - Wait-free synchronization [TOPLAS'91]: <https://dl.acm.org/citation.cfm?id=102808>

Back up slides