

Zookeeper AtomBroadcast

Emmanouil (Manos) Giortamis

Chair of Decentralized Systems Engineering (DSE)

Department of Computer Science

<https://dse.in.tum.de/>



ZAB: High-performance broadcast for
primary-backup systems

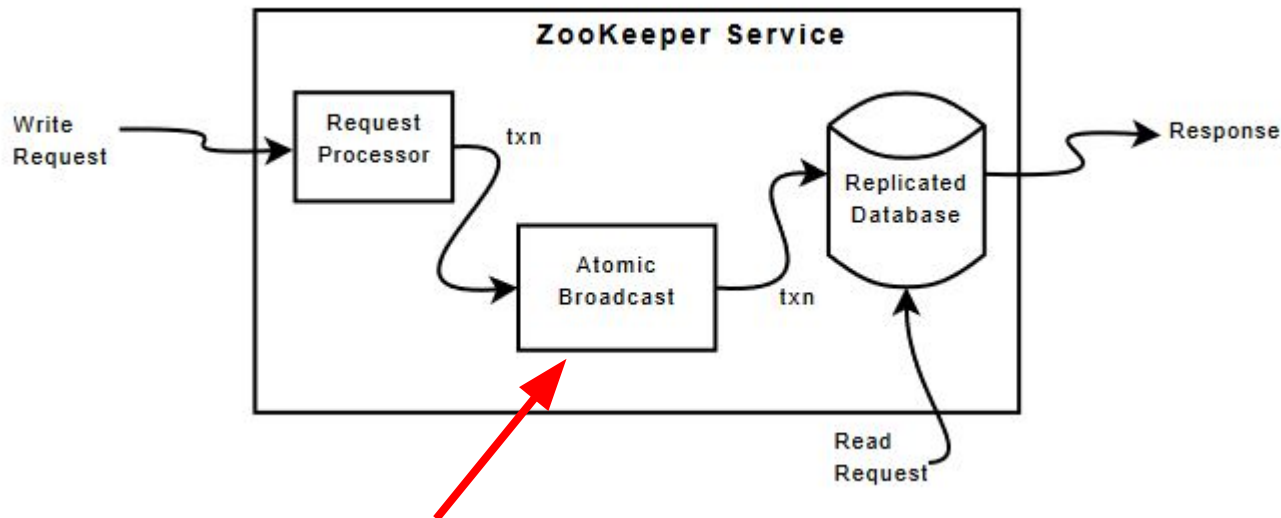
The problem

- Zookeeper uses a primary-backup scheme for replica consistency
- Primary nodes apply *incremental* and *idempotent* changes
- Application order must follow delivery order for state changes
 - “FIFO” order
 - *At-least-once* semantics
- Primary nodes can crash
 - Half-completed broadcasts
 - Fast and consistent recovery needed

Existing solutions did not satisfy the above requirements (atomic, custom ordering, fast recovery)

Solution: ZAB

- A crash-recovery Atomic Broadcast protocol
- Guarantees *at-least-once* semantics
- Implements custom “*Primary order*” ordering



- Useful primitive with many applications in distributed systems
- Atomic: either completes as a whole, or not at all
- Broadcast: transfer of messages to all recipients simultaneously
- Ordering: how messages from different broadcasts are ordered
- Various ordering semantics:
 - Total order
 - FIFO order
 - Causal order
 - ...

- Primary-backup replica roles
- Any replica can act as a primary
 - At most one active at a time
- Works with epochs
 - That change with each new primary
- State changes are called transactions (TXs)
 - Invoked only by a primary
- TXs are identified by an epoch and a counter pair (*zxid*)
 - In some epoch a new TX increments the counter
- Ordering as follows:
 - *zxid1* < *zxid2* iff:
 - *zxid1*.epoch < *zxid2*.epoch
 - *zxid1*.epoch = *zxid2*.epoch and *zxid1*.counter < *zxid2*.counter

- Uses TCP that preserves FIFO order
- (Safety) Properties:
 - Integrity: a process receives a message iff it was sent by some process
 - Agreement: any two processes deliver the same messages
 - Primary Order (PO):
 - If a primary broadcasts m before m' , then a recipient delivers m before m'
 - If two primaries broadcast two messages in different epochs, a recipient delivers them in the epoch order
 - Primary Integrity: a primary broadcasts iff it has delivered the TXs of previous epochs

- Two roles: *leader* and *follower*
- ZAB assumes a leader election to happen first
- All replicas know which is the new leader via an *oracle*
- It consists of three phases:
 - Discovery
 - Synchronization
 - Broadcast
- A leader is not “officially” leading until the synchronization phase completes
- A leader executes the steps of the follower as well
- A leader receives all changes from previous leaders (epochs) before broadcasting their own

Discovery phase

- Follows after a leader election algorithm
- Q: quorum

Leader l	Follower
	Send last epoch number e to l
Upon receiving messages from Q, propose new epoch e' s.t. $e' > e$ for all messages received in Q	
	Upon receiving e' , if $e' > e$, then $e = e'$. Reply with <i>ACK</i> with the highest <i>zxid</i>
Upon receiving <i>ACKs</i> from Q, select highest <i>zxid</i> . Receive missing TXs from followers (s.t. $zxid' < zxid$)	

Synchronization phase

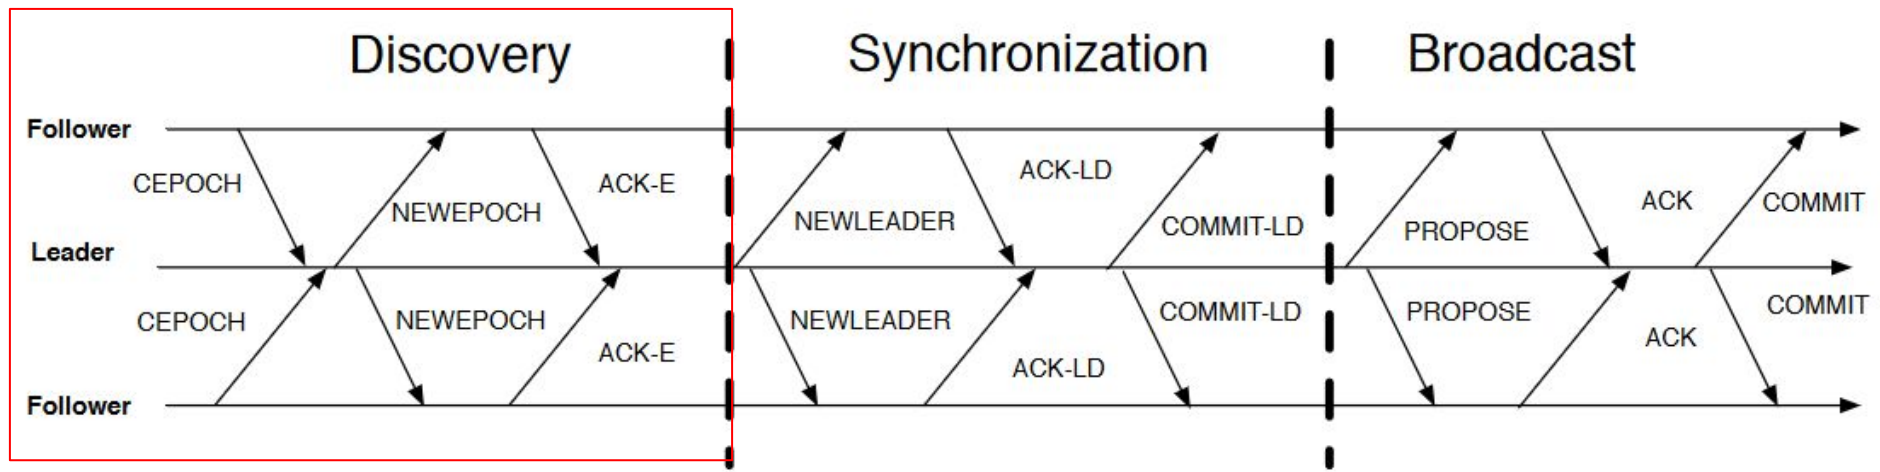
Leader <i>I</i>	Follower
Propose <i>NEWLEADER</i> and send highest <i>zxid</i> selected.	
	Set <i>I</i> as the new leader. Accept TXs with <i>zxid'</i> < <i>zxid</i> . Reply with <i>ACK</i>
Upon receiving <i>ACKs</i> from <i>Q</i> , send <i>COMMIT</i> message to all followers	
	Upon receiving <i>COMMIT</i> , deliver the previously accepted TXs

Broadcast phase

Leader	Follower
Increment counter and propose TXs Now $zxid' > \text{previous largest } zxid$	
	Accept the TXs and reply with <i>ACK</i>
Upon receiving <i>ACKs</i> from <i>Q</i> , send <i>COMMIT</i> to all followers	
	Upon receiving <i>COMMIT</i> , deliver the TXs

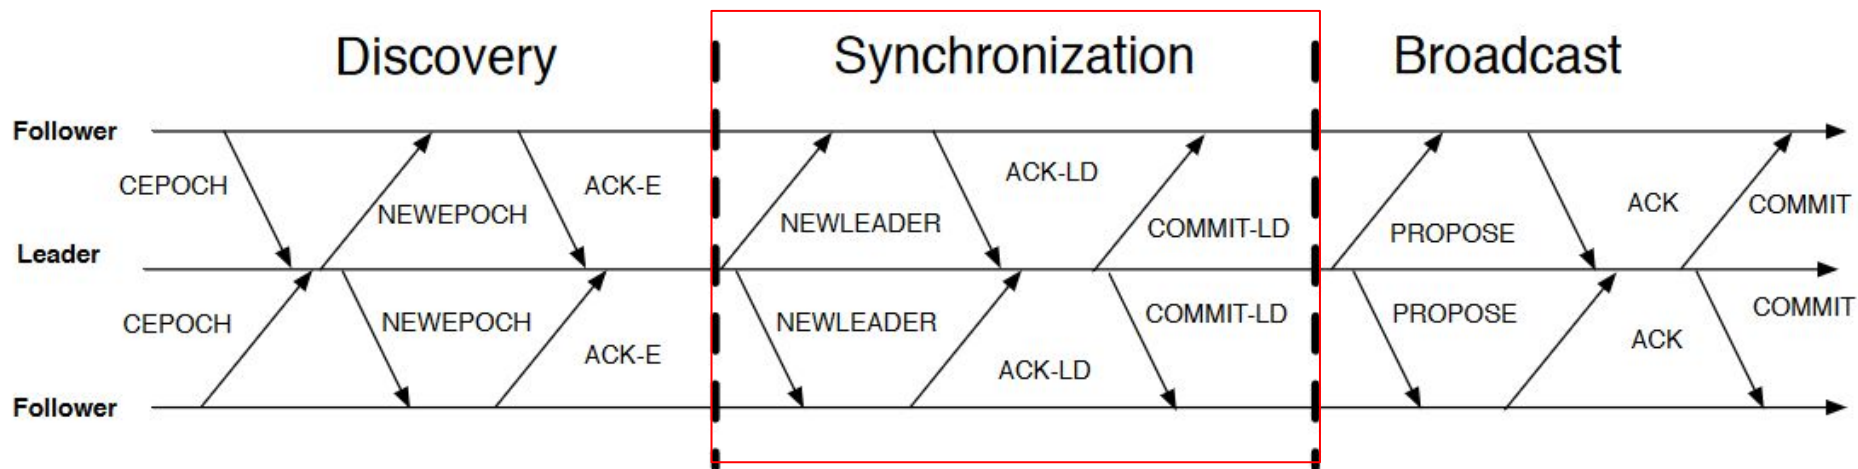
Complete picture

- CEPOCH: last epoch number
- NEWEPOCH: new epoch proposal
- ACK-E: acknowledgement of epoch proposal



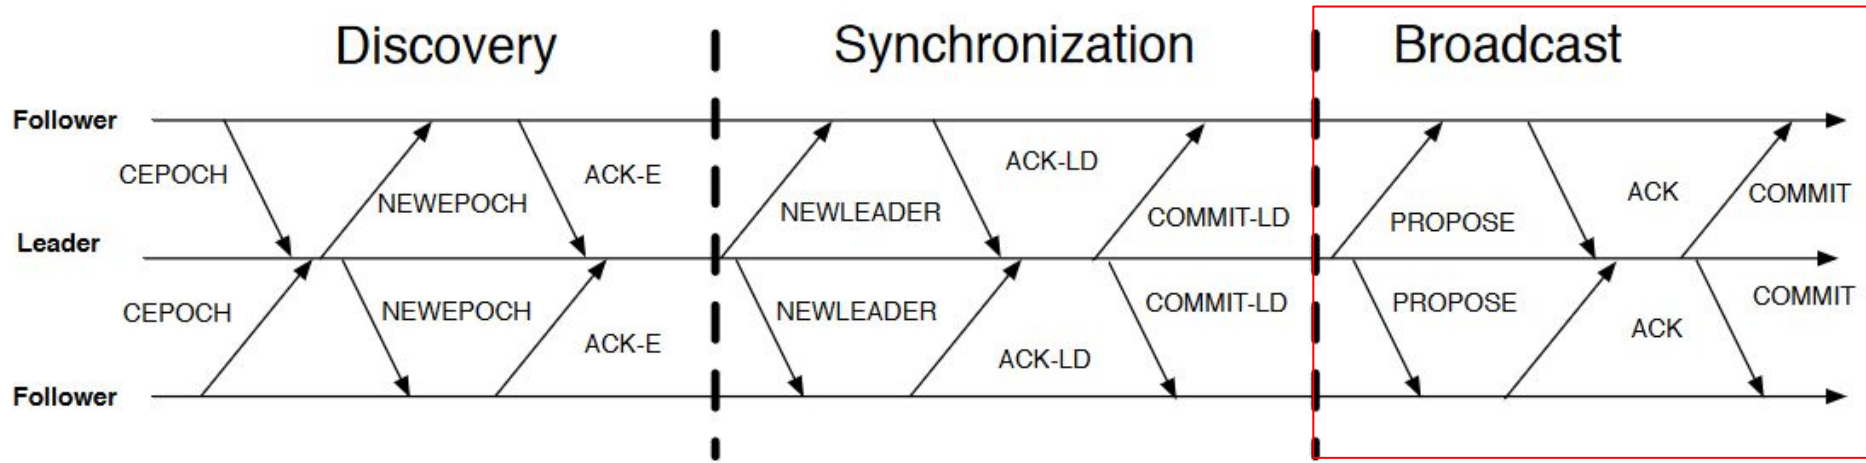
Complete picture

- NEWLEADER: propose self as new leader
- ACK-LD: acknowledgement of new leader proposal
- COMMIT-LD: commit leader proposal



Complete picture

- PROPOSE: propose new TX
- ACK: acknowledgement of leader's proposal
- COMMIT: leader commits proposal



- Leaders and followers use heartbeat messages
- A leader:
 - Receives heartbeats from Q
 - If not enough heartbeats, steps down
 - Triggers leader election
 - Moves to phase 1 (Discovery)
- A follower:
 - Follows a leader as long as heartbeat messages are received
 - If not, it abandons them
 - Triggers leader election
 - Moves to phase 1 (Discovery)

References

- ZAB: <https://marcoserafini.github.io/papers/zab.pdf>
- Zookeeper: https://www.usenix.org/legacy/event/atc10/tech/full_papers/Hunt.pdf